

# Module IN3031 / INM378 Digital Signal Processing and Audio Programming

Tillman Weyde

[t.e.weyde@city.ac.uk](mailto:t.e.weyde@city.ac.uk)



# Image Processing Texts

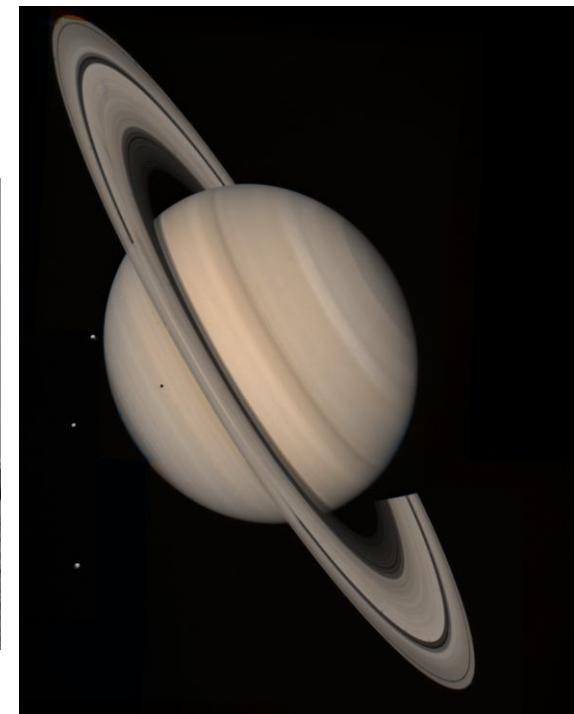
John C. Russ: The image processing handbook.  
IEEE Press, 3rd edition, 1999.

Rafael C. Gonzalez and Richard E. Woods:  
Digital image processing. Addison Wesley, 1993.  
Chapters 1-2 available at: <http://www.imageprocessingplace.com>



# What is an image? (1)

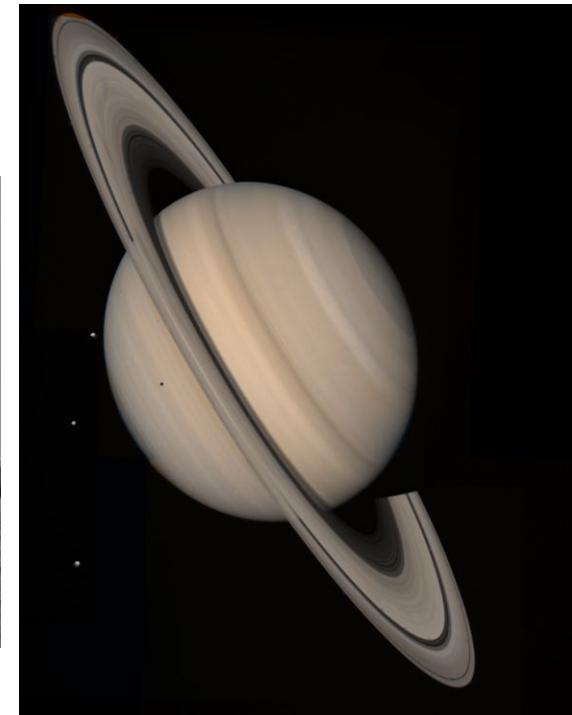
- A visible impression obtained by a camera, telescope, microscope, or other device, or displayed on a computer or video screen - *Oxford Dictionaries*





# What is an image? (2)

- In form of function  $f(x,y)$
- Most images are rectangular



# What is Digital Image Processing?

- It concerns the transformation of an image to a digital format and its processing by computers
- **Processing:** analysis, synthesis, modification
- Why process?
  - Facilitate storage and transmission
  - Restore or enhance images
  - Extract information from images
  - ...

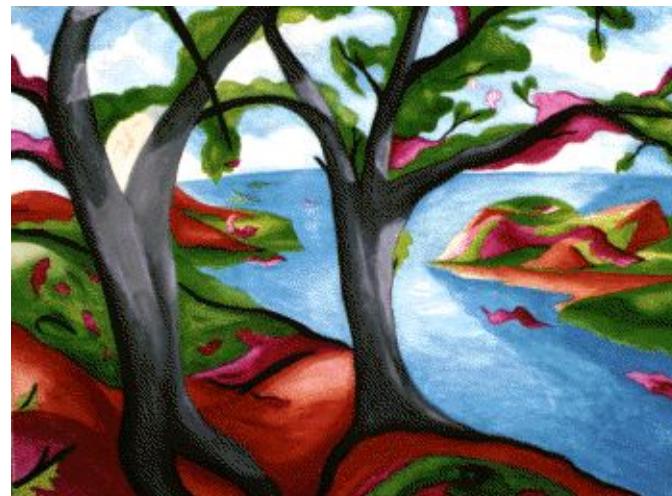


# What is a digital image?

Discrete samples  $f[x,y]$  representing the continuous image  $f(x,y)$

Image size:  $\max(x) \times \max(y)$ , e.g. 800x600

**Pixels** (“picture elements”): elements of the 2-D array





# Topics in Digital Image Processing

## Image compression



PNG (300 kB)



JPG (3 kB)



# Topics in Digital Image Processing

## Medical Imaging





# Topics in Digital Image Processing

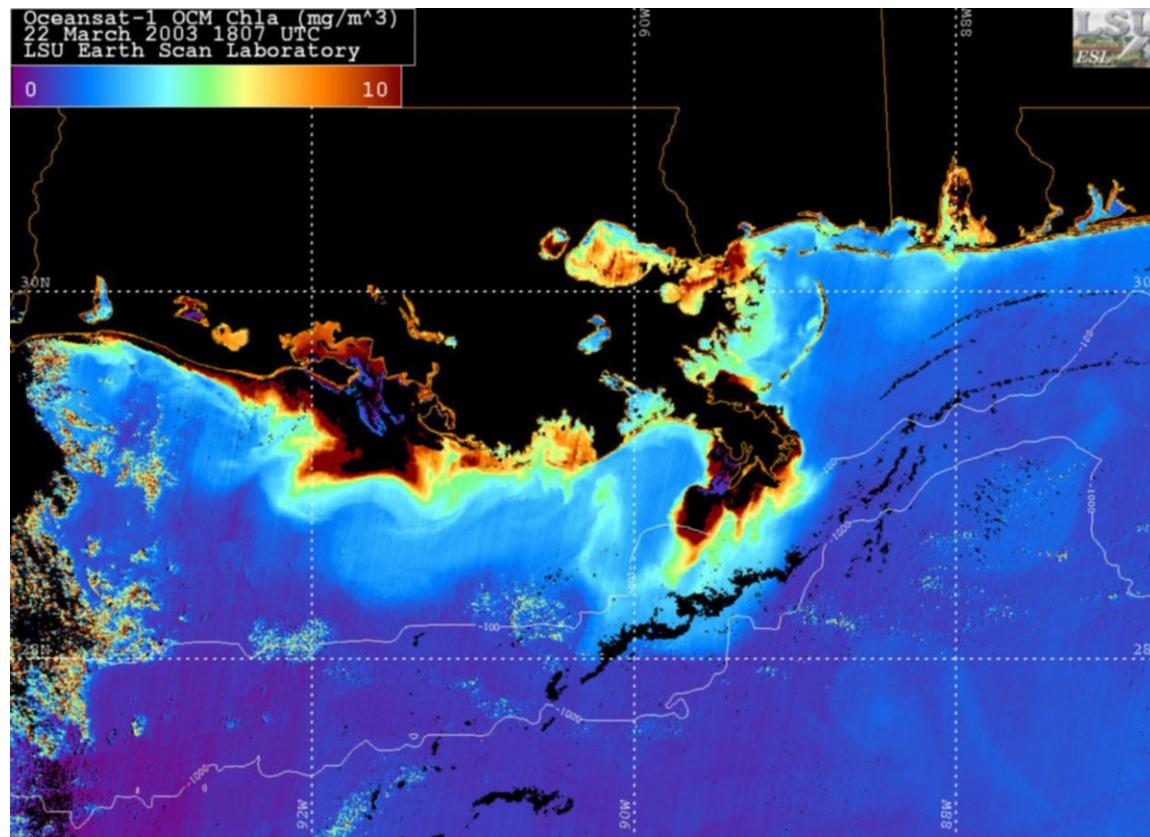
## Digital Image Restoration/Enhancement





# Topics in Digital Image Processing

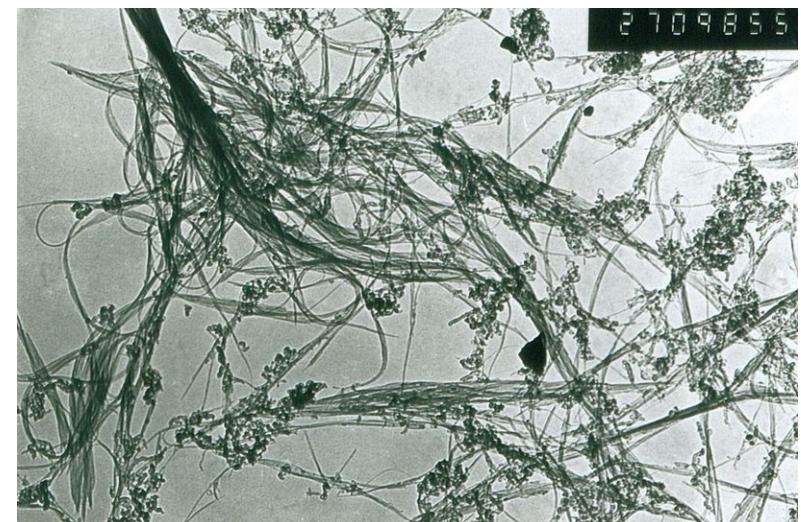
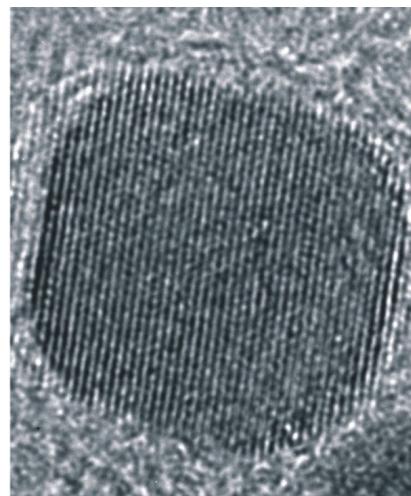
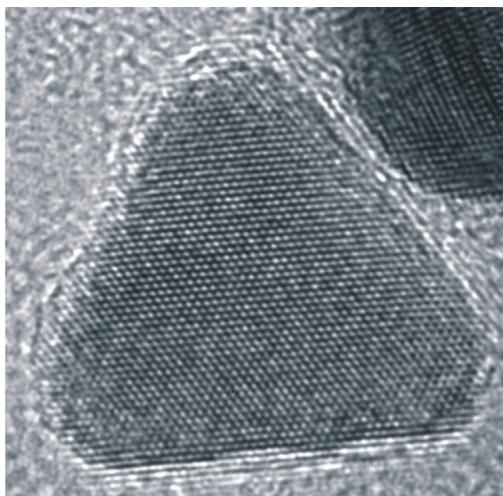
## Remote Sensing





# Topics in Digital Image Processing

## Microscopy



Transmission electron microscopy (TEM) images



# Topics in Digital Image Processing

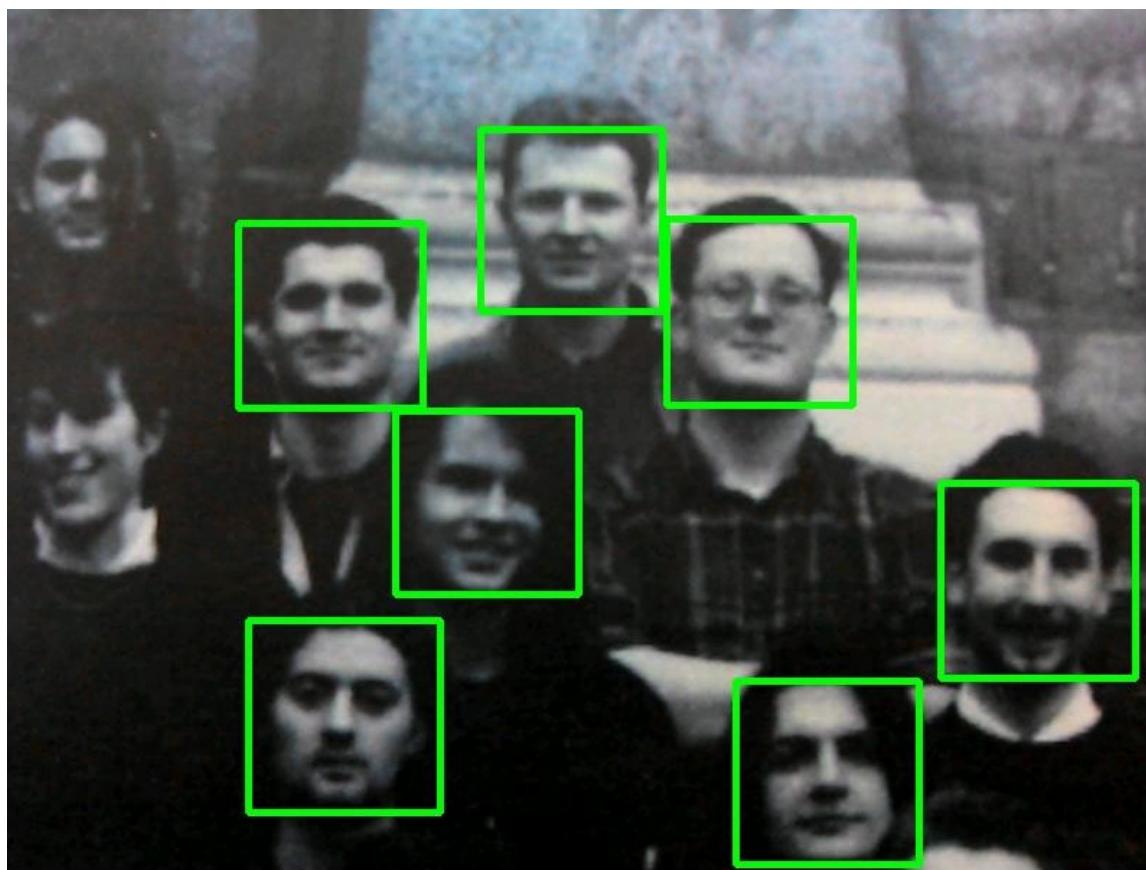
## Image Registration





# Topics in Digital Image Processing

## Face Detection





# Topics in Digital Image Processing

## Handwriting recognition

InkOverlay Example

Erase

Winter is here. Go to the store and buy some snow shovels.

Winter is here. Go to the store and buy some snow shovels.

$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-i2\pi \frac{k}{N} n}$

MathJax

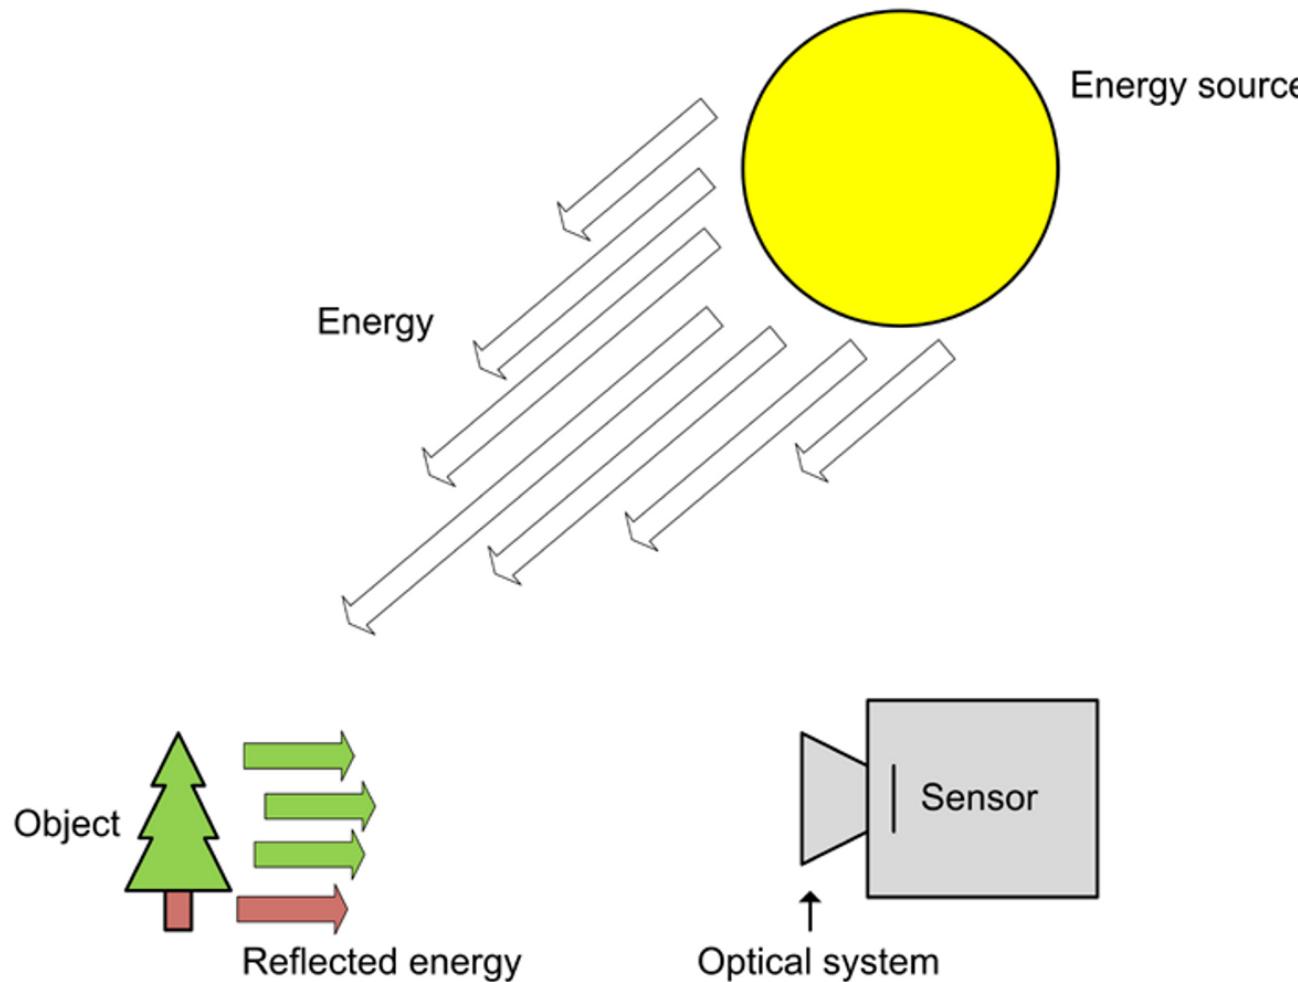
LaTeX MathML

$X_k = \sum_{n=0}^{N-1} x_n \cdot e^{-i2\pi \frac{k}{N} n}$

# Digital Image Acquisition



# Digital Image Acquisition (1)



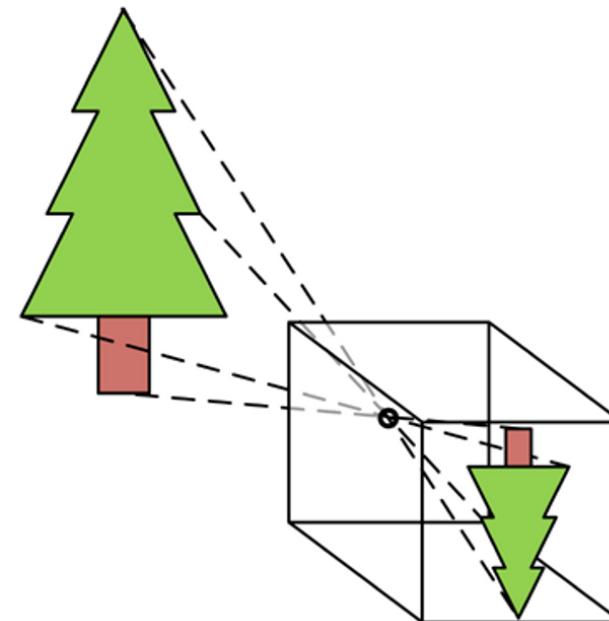
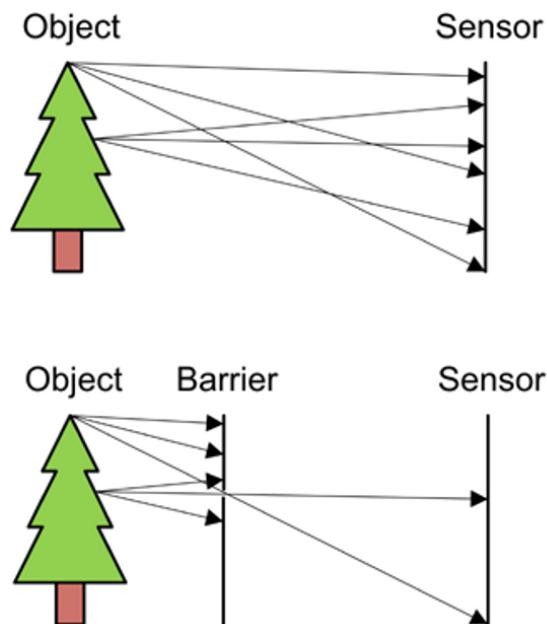
Overview of the typical image acquisition process, with the sun as light source, a tree as object and a camera to capture the image.

An analog camera would use a film where the digital camera uses a sensor.



# Digital Image Acquisition (2)

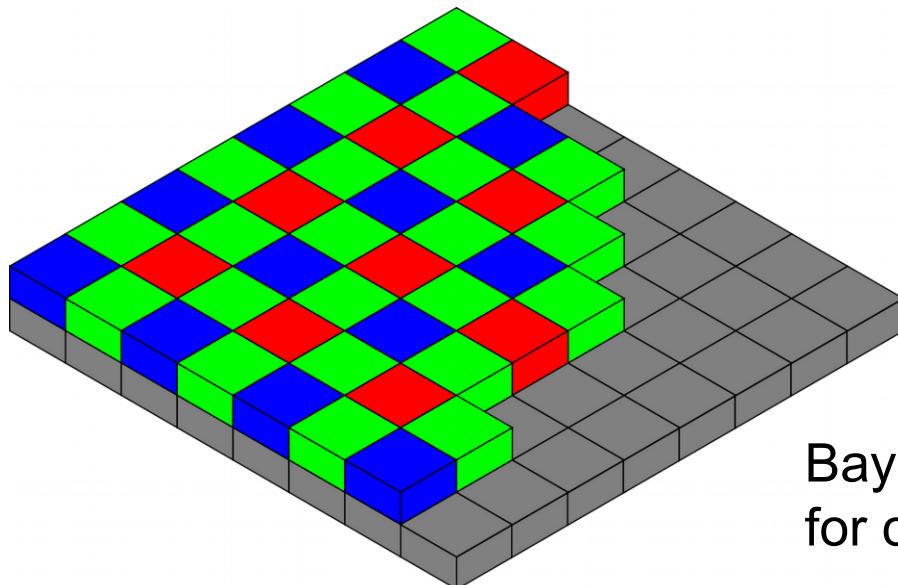
- Light reflected from the object has to be captured
- A barrier (also called **pinhole model**) is placed between the object of interest and the sensing material





# Digital Image Acquisition (3)

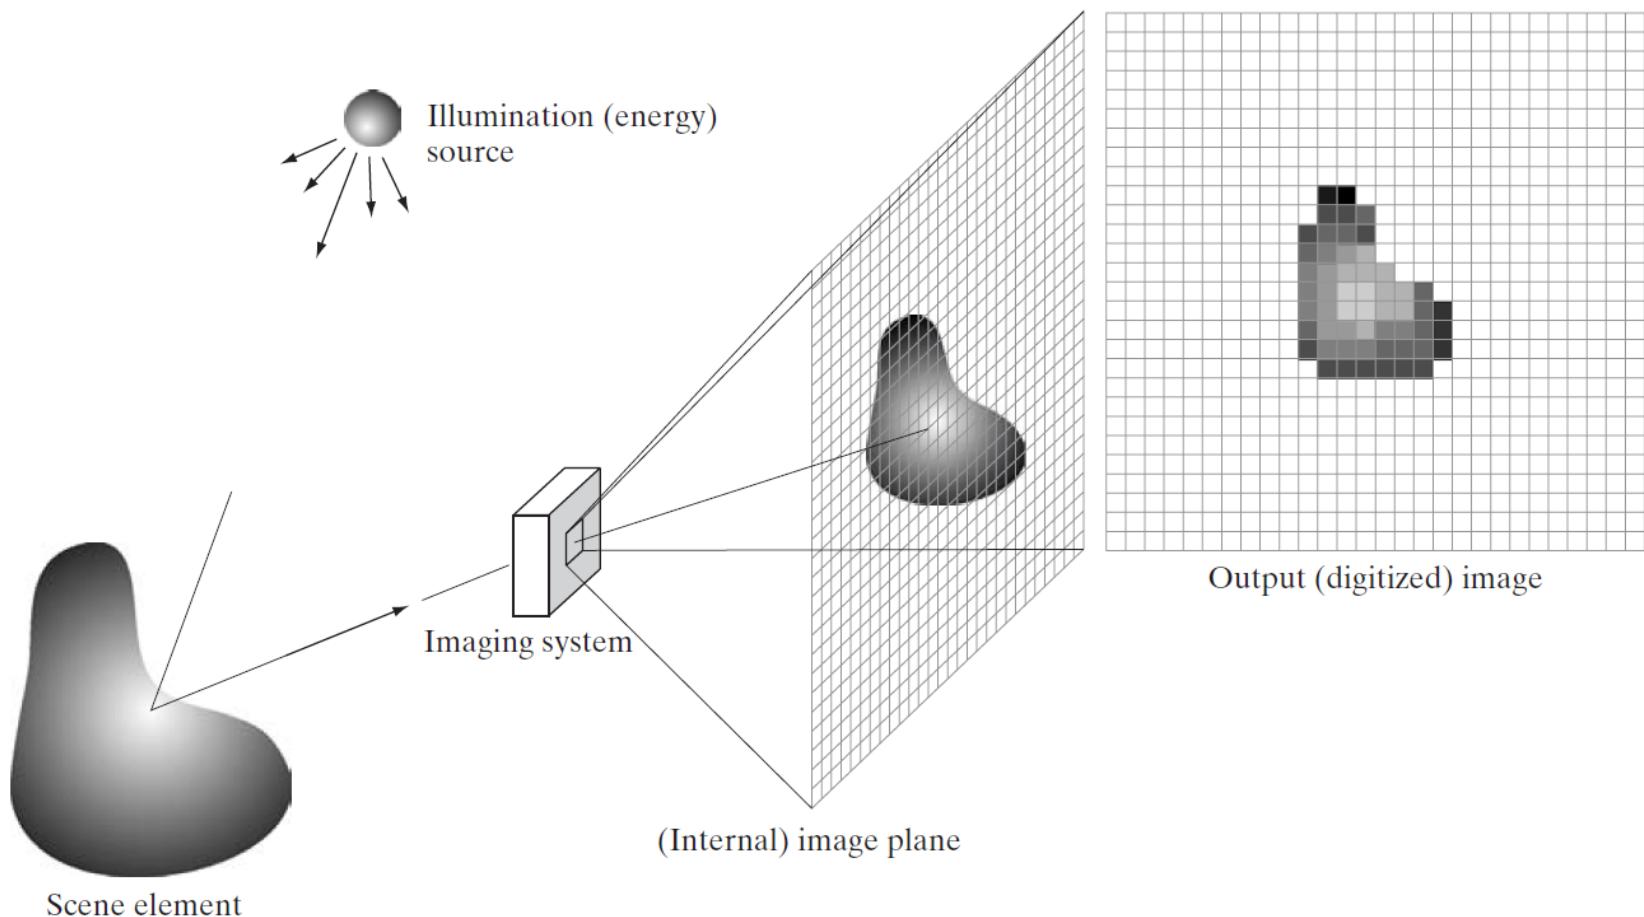
- Image sensors are used for recording the input light
- An image sensor consists of a 2D array of cells
- Each cell measures light and converts it into voltage
- Digital image sensors: CCD (charge-coupled devices), CMOS (complementary metal–oxide–semiconductors)



Bayer arrangement on sensor  
for colour image acquisition



# Sampling and Quantization (1)

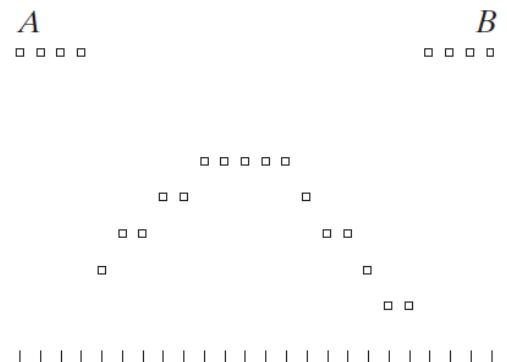
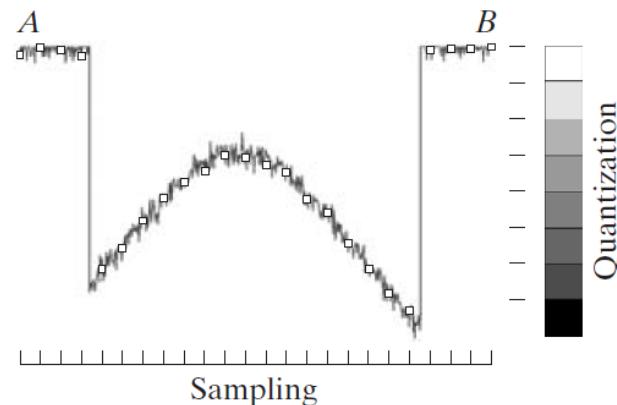
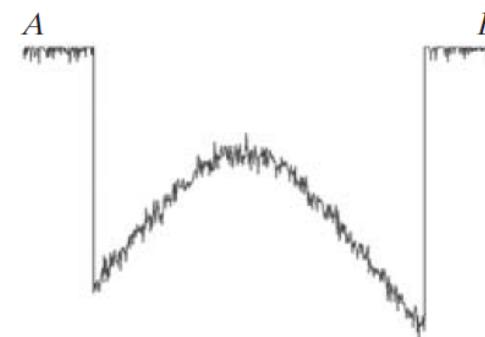
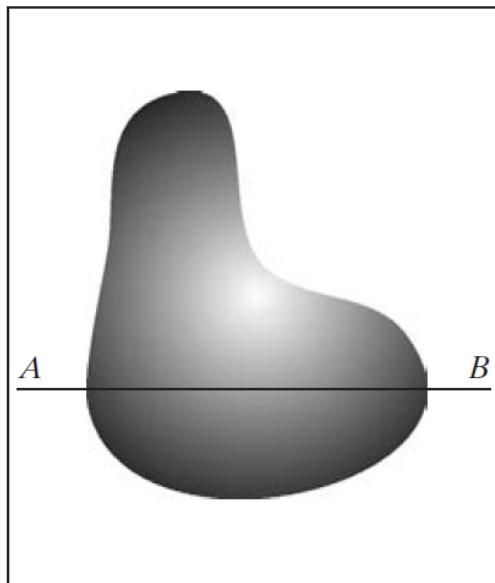


# Sampling and Quantization (2)

- To create a digital image, we need to convert the continuous sensed data into digital form. This involves two processes: **sampling and quantization**.
- To convert it to digital form, we have to sample the function  $f(x,y)$  in both coordinates and in amplitude.
- Digitizing the coordinate values is called **sampling**.
- Digitizing the amplitude values is called **quantization**.

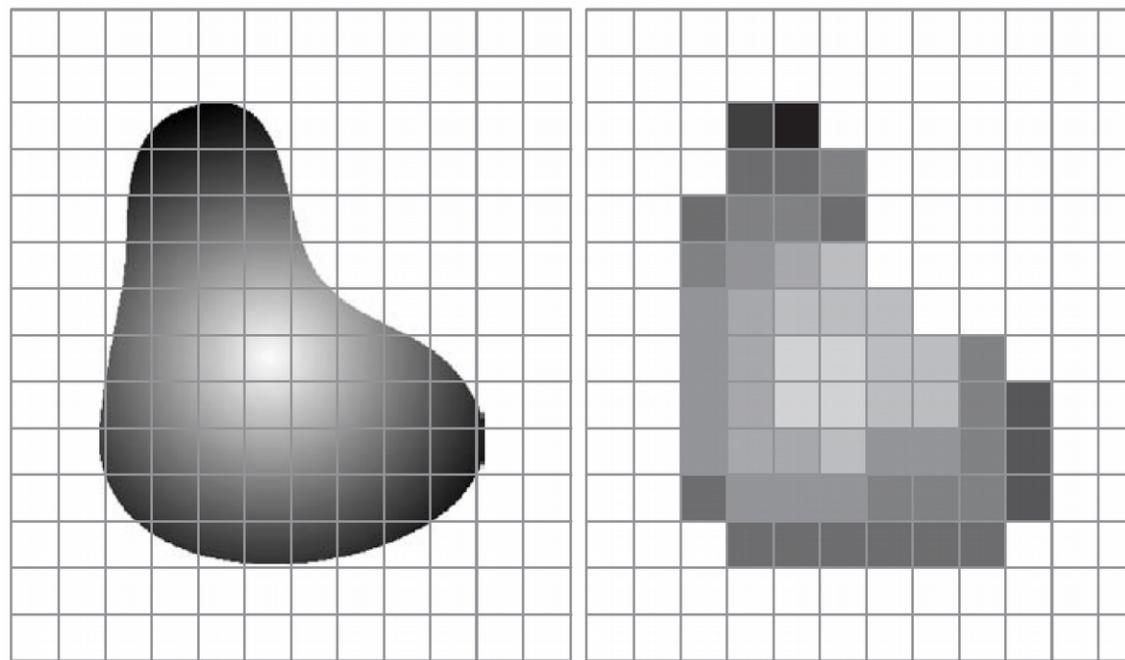


# Sampling and Quantization (3)



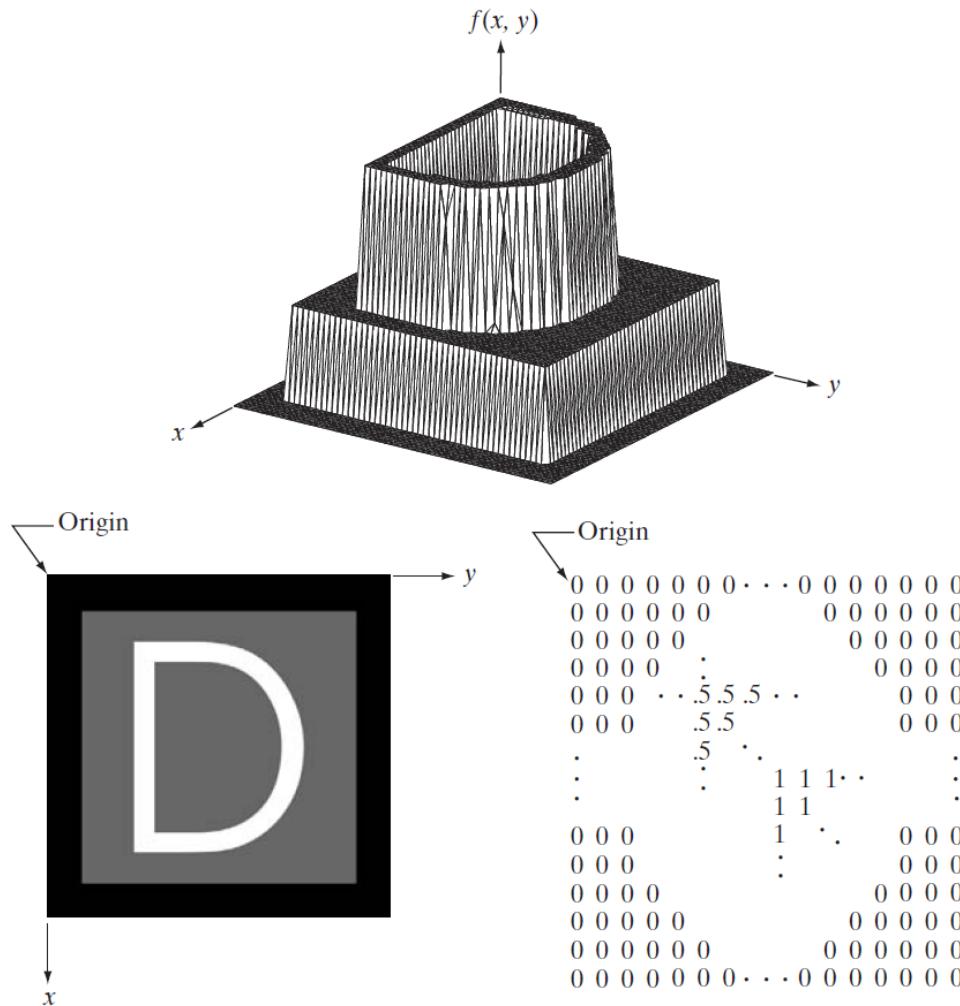
# Sampling and Quantization (4)

- The quality of a digital image is determined to a large degree by the number of samples and discrete intensity levels used in sampling and quantization.



# Representing Digital Images

# Image Representation (1)



# Image Representation (2)

- A digital image can be conveniently represented by a  $N \times M$  matrix:

$$A = \begin{bmatrix} a_{0,0} & a_{0,1} & \dots & a_{0,N-1} \\ a_{1,0} & a_{1,1} & \dots & a_{1,N-1} \\ \vdots & \vdots & \dots & \vdots \\ a_{M-1,0} & a_{M-1,1} & \dots & a_{M-1,N-1} \end{bmatrix}$$

Each element of the array is called a pixel.



# Spatial Resolution

- Sampling is the principal factor in determining the **spatial resolution** of an image:



256x256



128x128

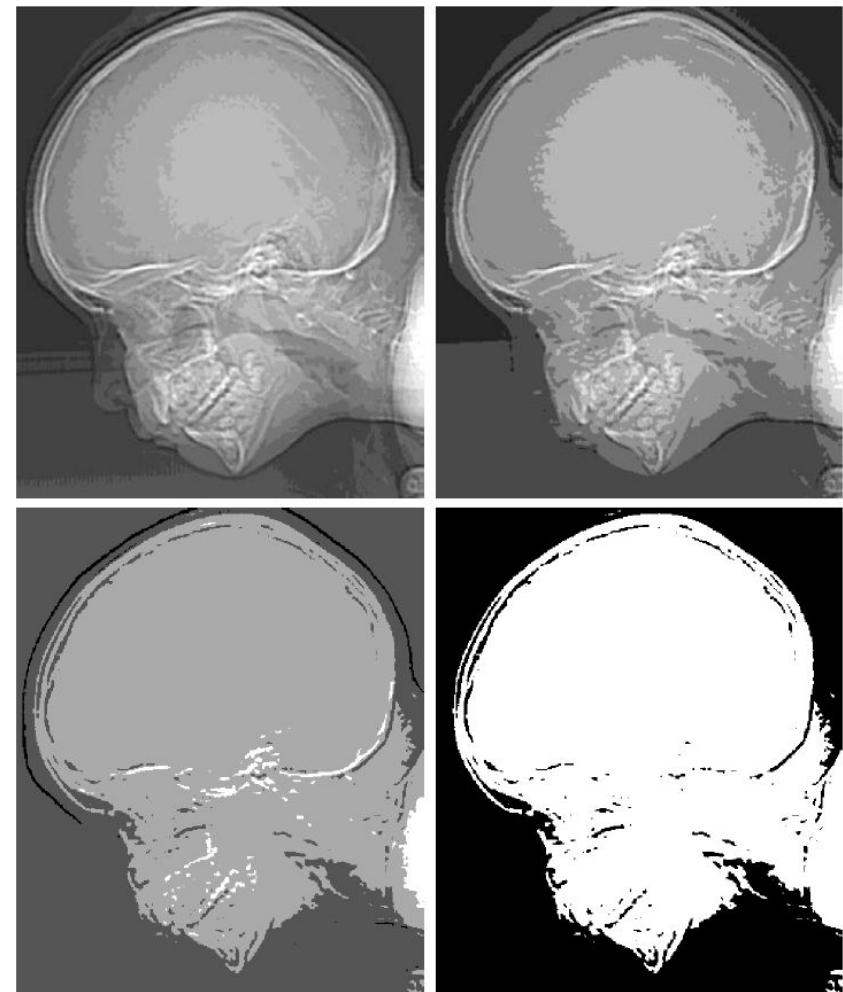


64x64



# Gray Levels

- **Gray-level resolution** refers to the smallest representable change in gray level.
- Usually an integer power of 2 (e.g. 8 bits = 256) as number of levels
- Figure: image in 16, 8, 4, and 2 gray levels





# Image Negatives

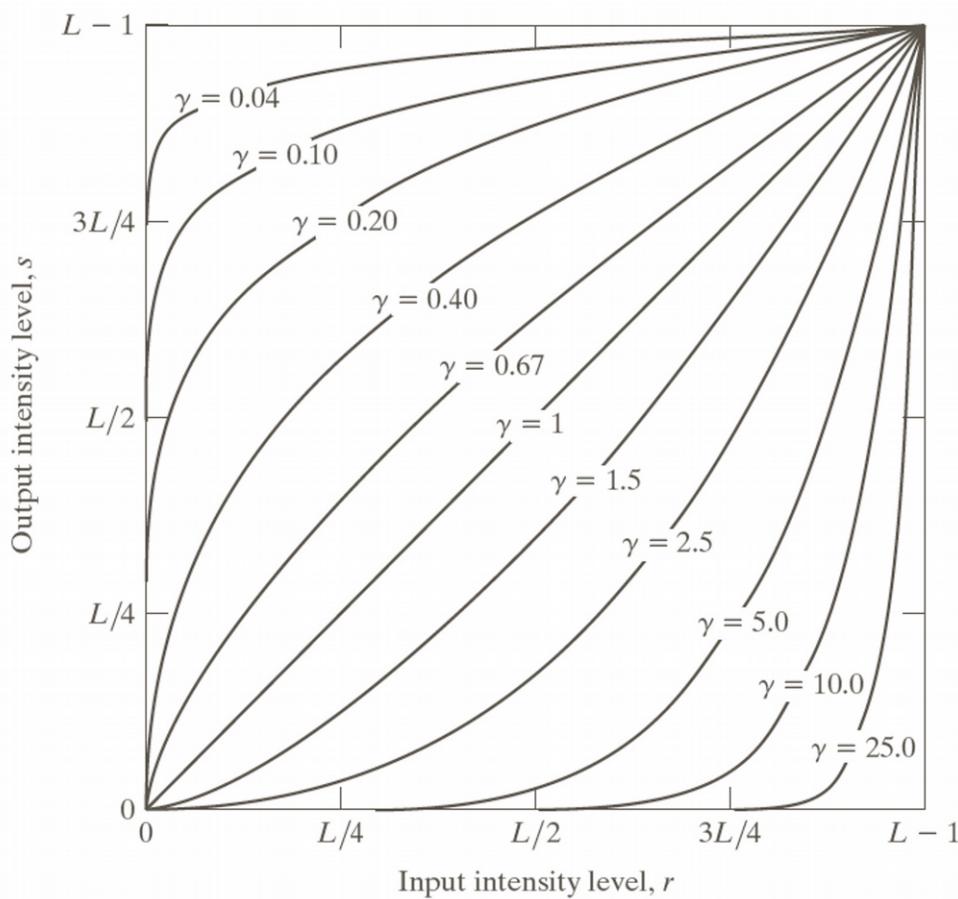
- Given an image with  $L$  gray levels, its negative transformation for pixel with level  $r$  is given by:  $s = L-1-r$





# Gamma Correction (1)

$$s = cr^\gamma \quad (c: \text{constant})$$



- In some devices (e.g. CRT TVs) the gamma curve is nonlinear, and needs to be corrected.
- Gamma correction is important in displaying images correctly (otherwise they might be either bleached or dark).



# Gamma Correction (2)



$\gamma = 0.5$



$\gamma = 1.0$



$\gamma = 1.5$

# Elementary Image Processing Operations

- Addition:  $C[i][j] = A[i][j]+B[i][j]$
- Subtraction:  $C[i][j] = A[i][j]-B[i][j]$
- Multiplication by a constant  $c$ :  $C[i][j] = c A[i][j]$

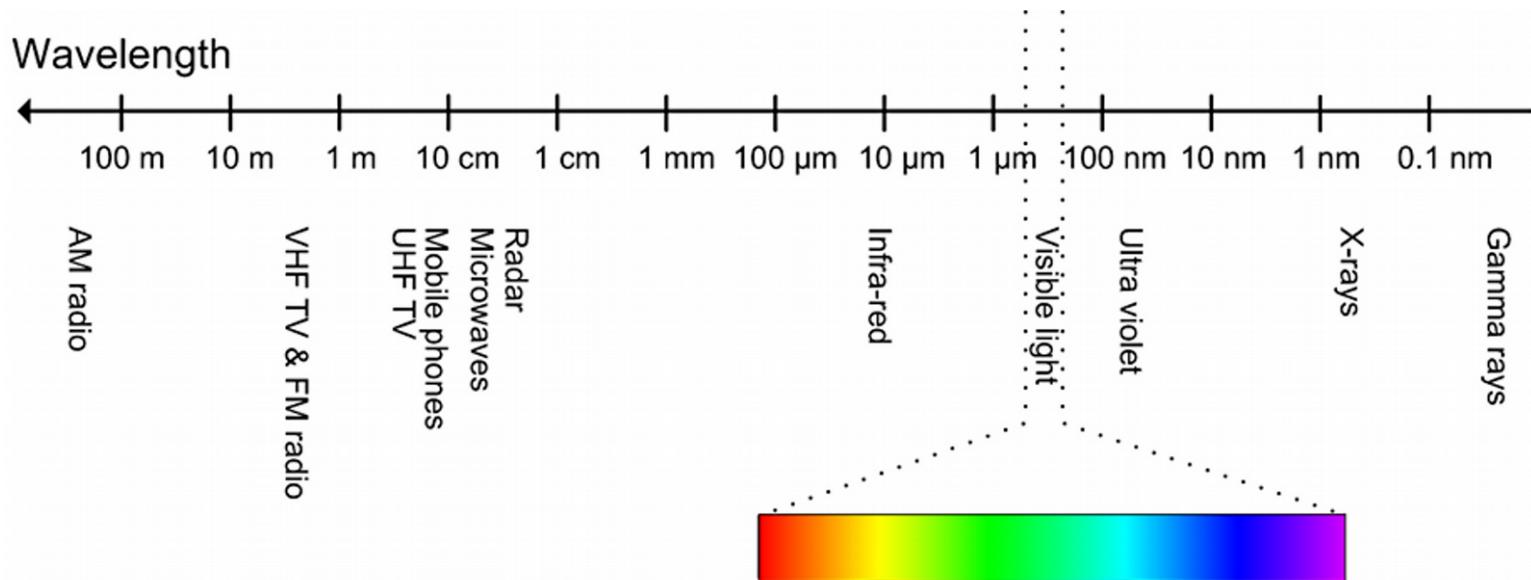


# Colour Image Processing



# Colour Image Processing (1)

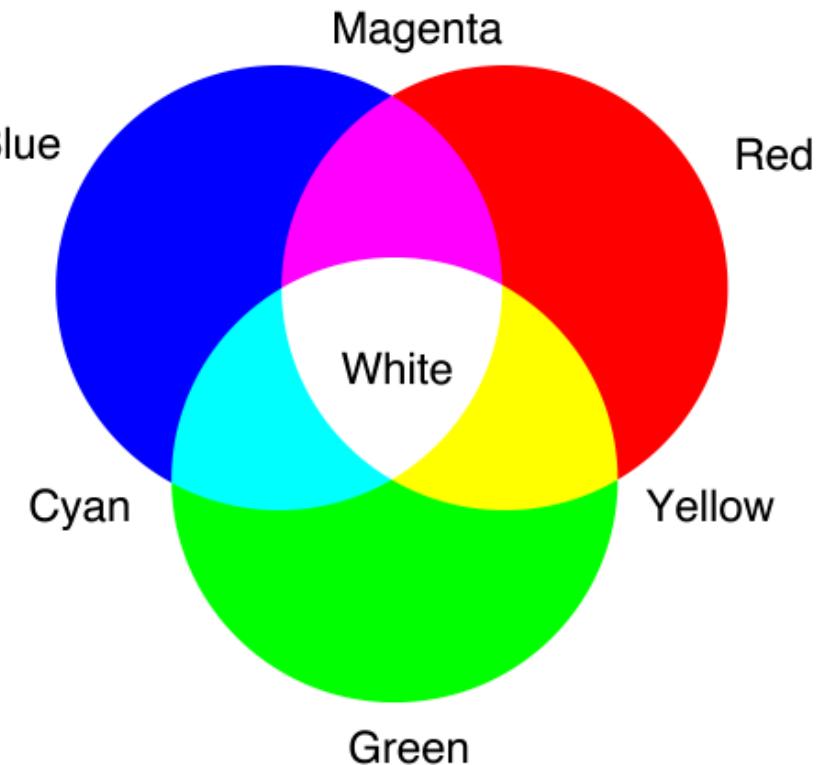
- Visible light is composed of a relatively narrow band of frequencies in the electromagnetic spectrum.
- Chromatic light spans the electromagnetic spectrum from approximately 400 to 700 nm.





# Colour Image Processing (2)

- 6 to 7 million cones in the human eye can be divided into 3 principal sensing categories, corresponding roughly to red (R), green (G), and blue (B) – also called **primary colours**.



- The primary colors can be added to produce the **secondary colors** of light – cyan (C), magenta (M), and yellow (Y).



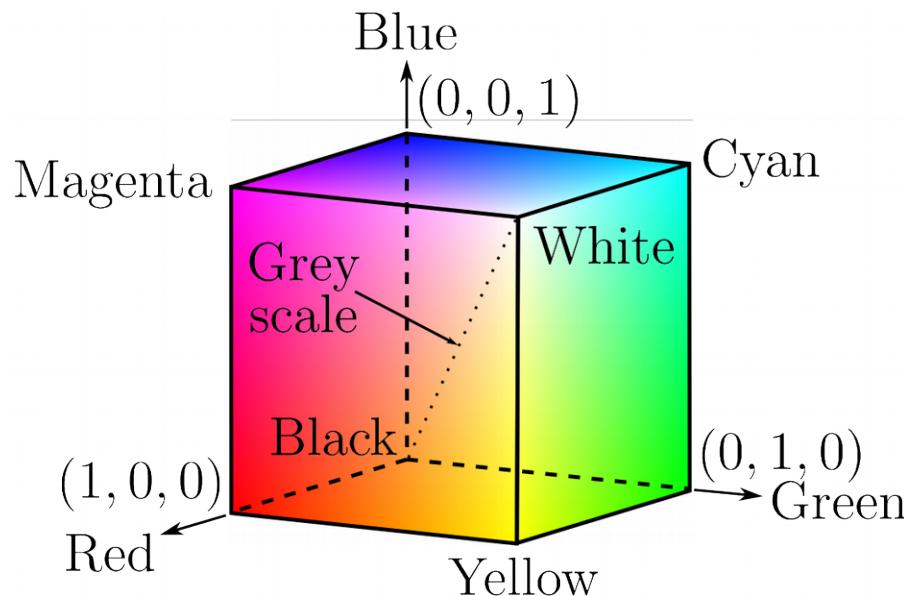
# Colour Image Processing (3)

- Characteristics used to distinguish one color from another are brightness, hue, and saturation.
  - **Hue** represents the dominant color as perceived by an observer.
  - **Saturation** refers to the relative purity or the amount of white light mixed with a hue.
  - **Brightness** embodies the notion of **intensity**.



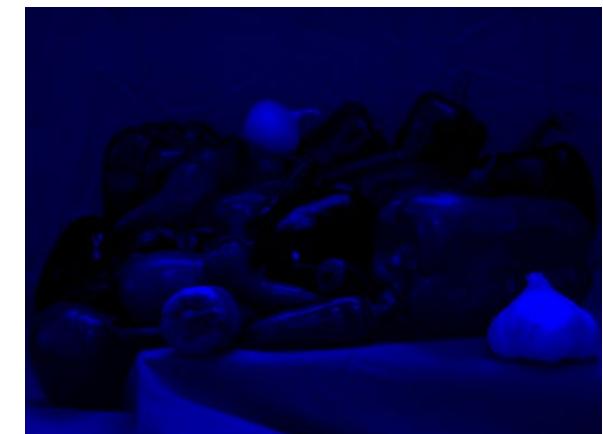
# Colour Models (1)

- The **RGB** (red, green, blue) Colour Model: used in colour monitors, most video cameras and many file formats.
- Red, green, and blue are each 8-bit images.





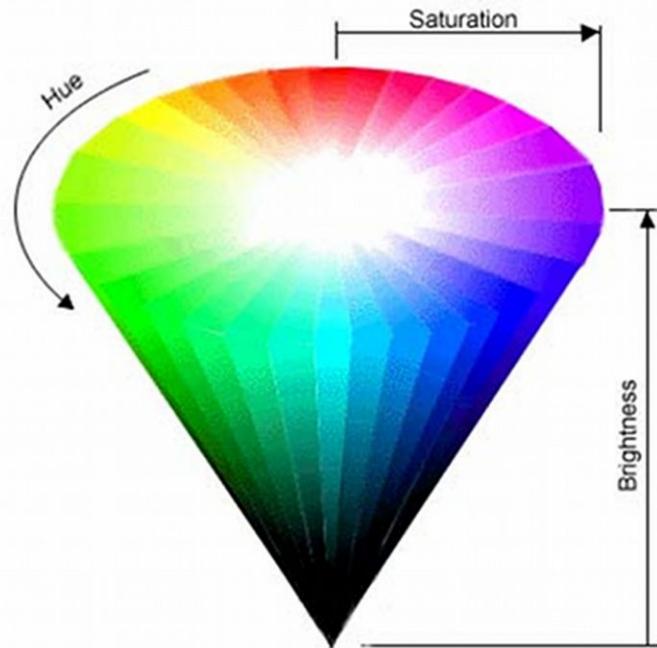
# Colour Models (2)





# Colour Models (3)

- **CMY** (cyan, magenta, yellow) and **CMYK** (cyan, magenta, yellow, black) models for color printing.
- **HSI** (hue, saturation, intensity) model, which corresponds closely with the way humans describe and interpret color.



# Colour Models (4)

- Converting RGB to CMY:

$$\begin{bmatrix} C \\ M \\ Y \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} - \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

# Colour Models (5)

- Converting RGB to HSI:

$$H = \begin{cases} \theta & \text{if } B \leq G \\ 360 - \theta & \text{if } B > G \end{cases}$$

$$\theta = \cos^{-1} \left\{ \frac{\frac{1}{2} [(R-G) + (R-B)]}{[(R-G)^2 + (R-B)(G-B)]^{1/2}} \right\}$$

$$S = 1 - \frac{3}{(R+G+B)} [\min(R, G, B)]$$

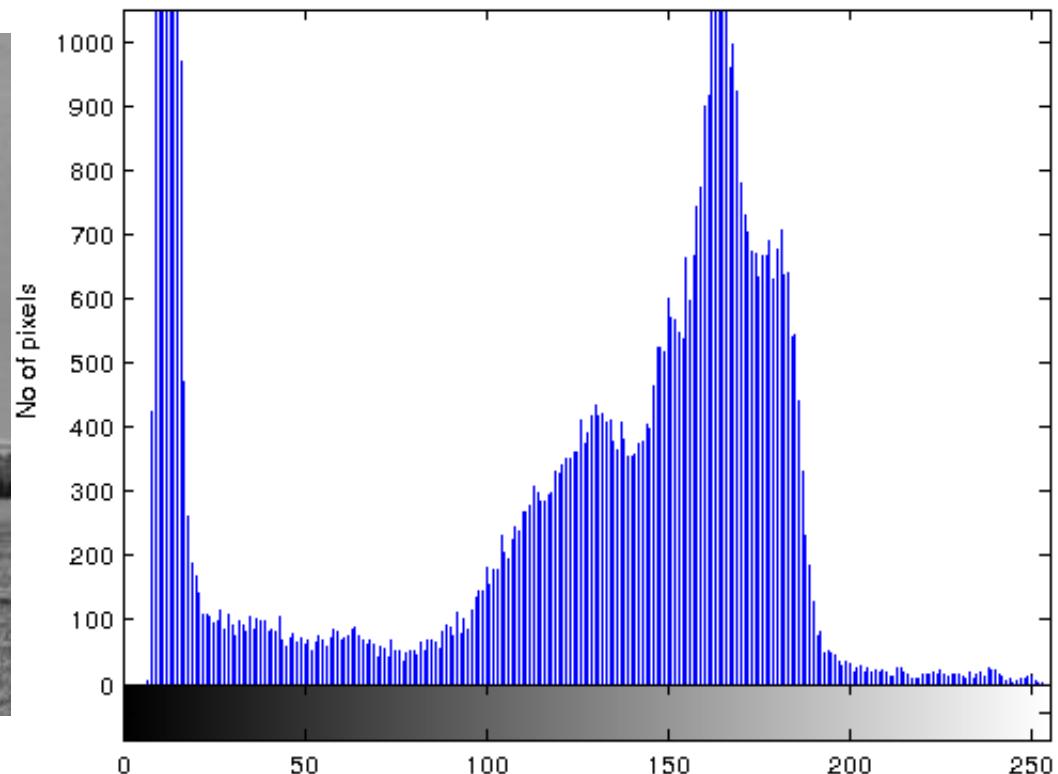
$$I = \frac{1}{3} (R+G+B)$$

# Histograms



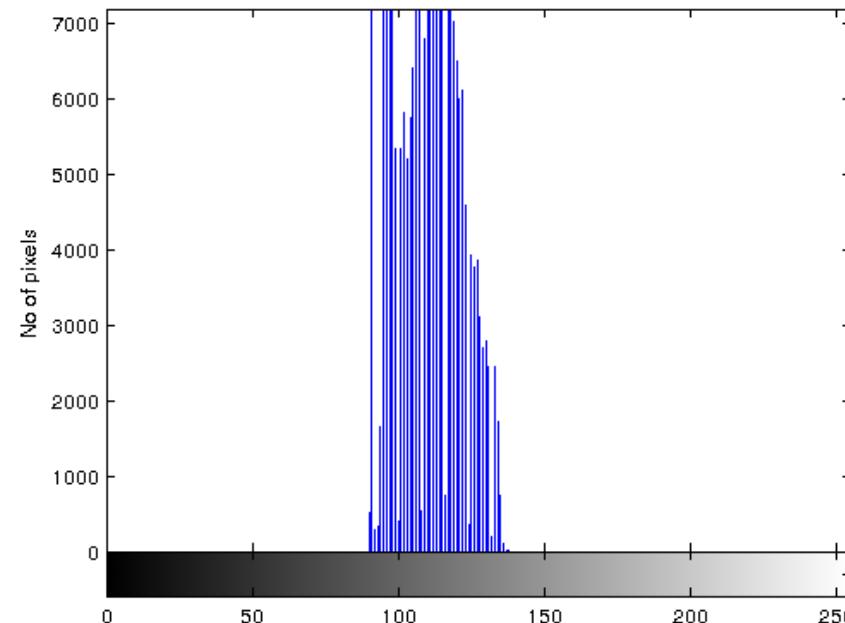
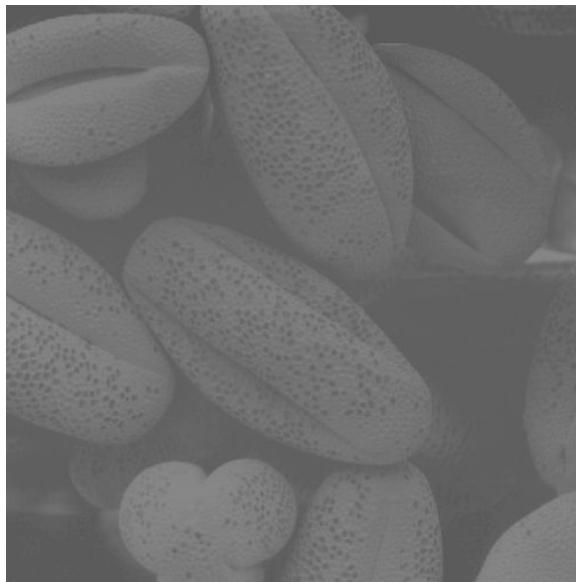
# Image Histograms (1)

- Distribution of the pixel intensity values



# Image Histograms (2)

- Histogram manipulation can be used for image enhancement.
- Figure: image with low **contrast**



# Histogram Equalisation (1)

- **Histogram equalisation:** method for contrast adjustment
- Increases the contrast of a low-contrast image
- **Goal:** find a transformation

$$s = T(r)$$

that is applied to each pixel of the input image  $x[m, n]$ , such that a uniform distribution of gray levels results for the output image  $y[m, n]$ .

# Histogram Equalisation (2)

- Implementation:
  - Compute the image histogram ( $n$ : total number of pixels,  $n_k$ : number of pixels with gray level  $r_k$ ):

$$p_r(r_k) = \frac{n_k}{n} \quad k=0,1,2,\dots,L-1$$

Compute the cumulative histogram:

$$s_k = \sum_{j=0}^k p_r(r_j)$$

New gray level values:

$$h(r_k) = \frac{s_k - s_{min}}{(MN) - s_{min}} (L-1)$$

# Histogram Equalisation (3)

- Example:

52	55	63	67	63
63	59	55	90	90
63	59	68	90	68
63	58	68	55	63
67	52	68	59	55

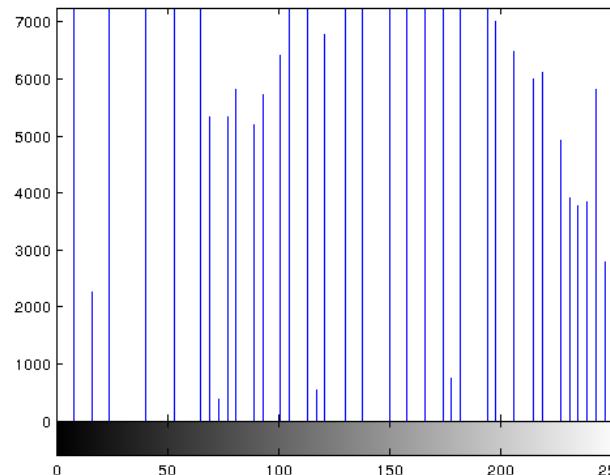
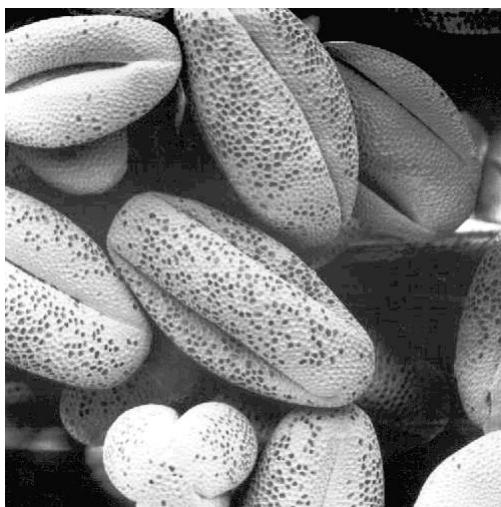
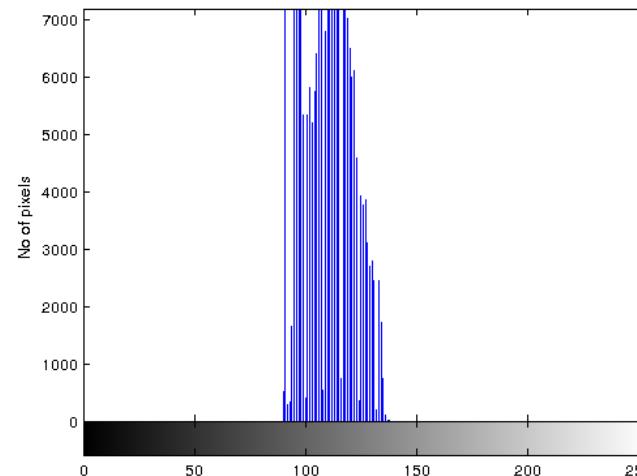
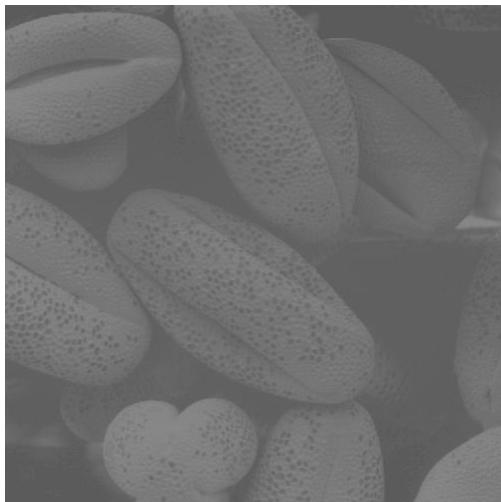
Image

52	2	2	0
55	4	6	44
58	1	7	55
59	3	10	89
63	6	16	155
67	2	18	177
68	4	22	222
90	3	25	255

Pixels, histogram, cumulative histogram, new pixel values



# Histogram Equalisation (4)



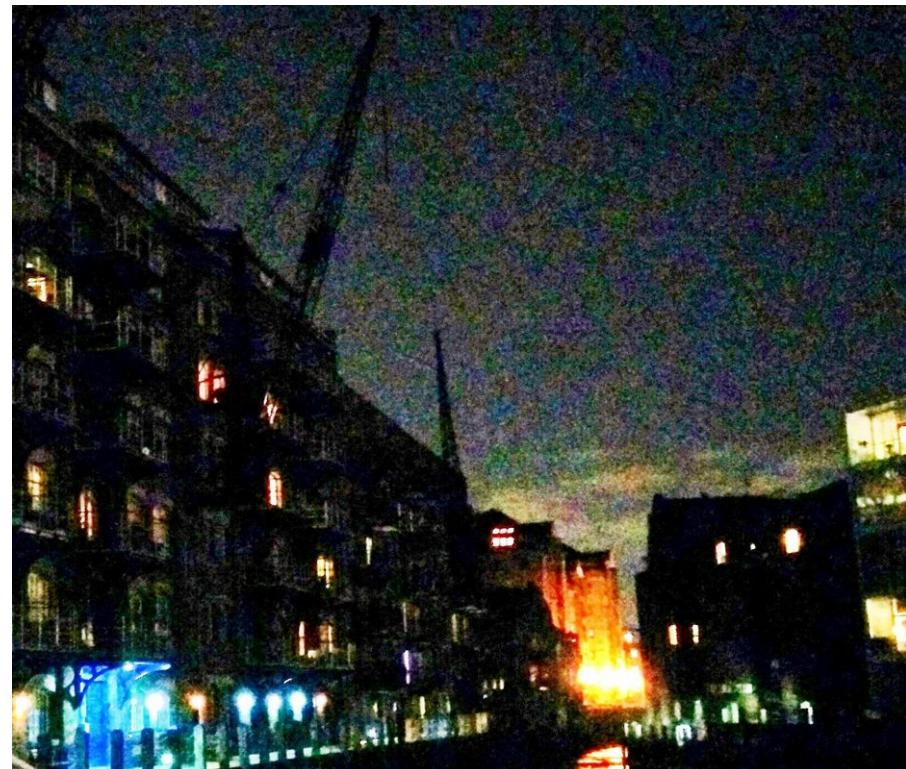
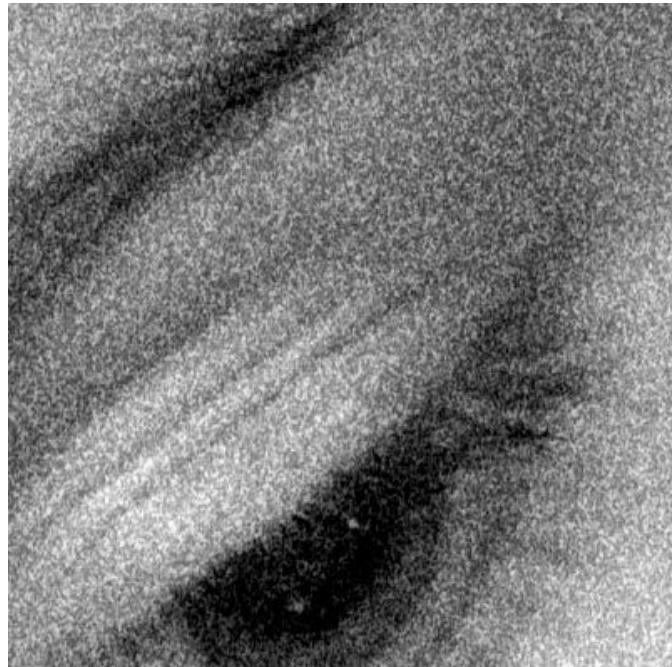


# Noise in Image Processing



# Noise in Image Processing (1)

- Digital images are corrupted by **noise** either during image acquisition (film grain noise, CCD noise...) or during image transmission.



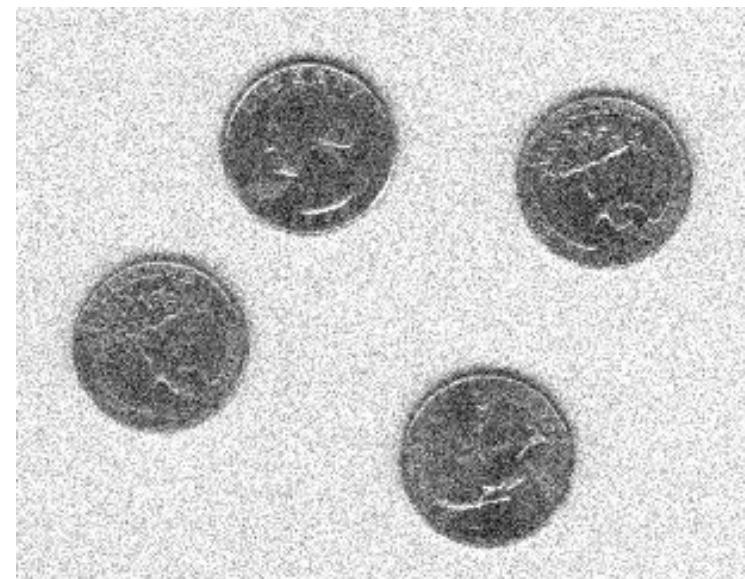


# Noise in Image Processing (2)

- Types of noise: salt & pepper, Gaussian, uniform...



Salt & Pepper Noise  
(impulsive noise)

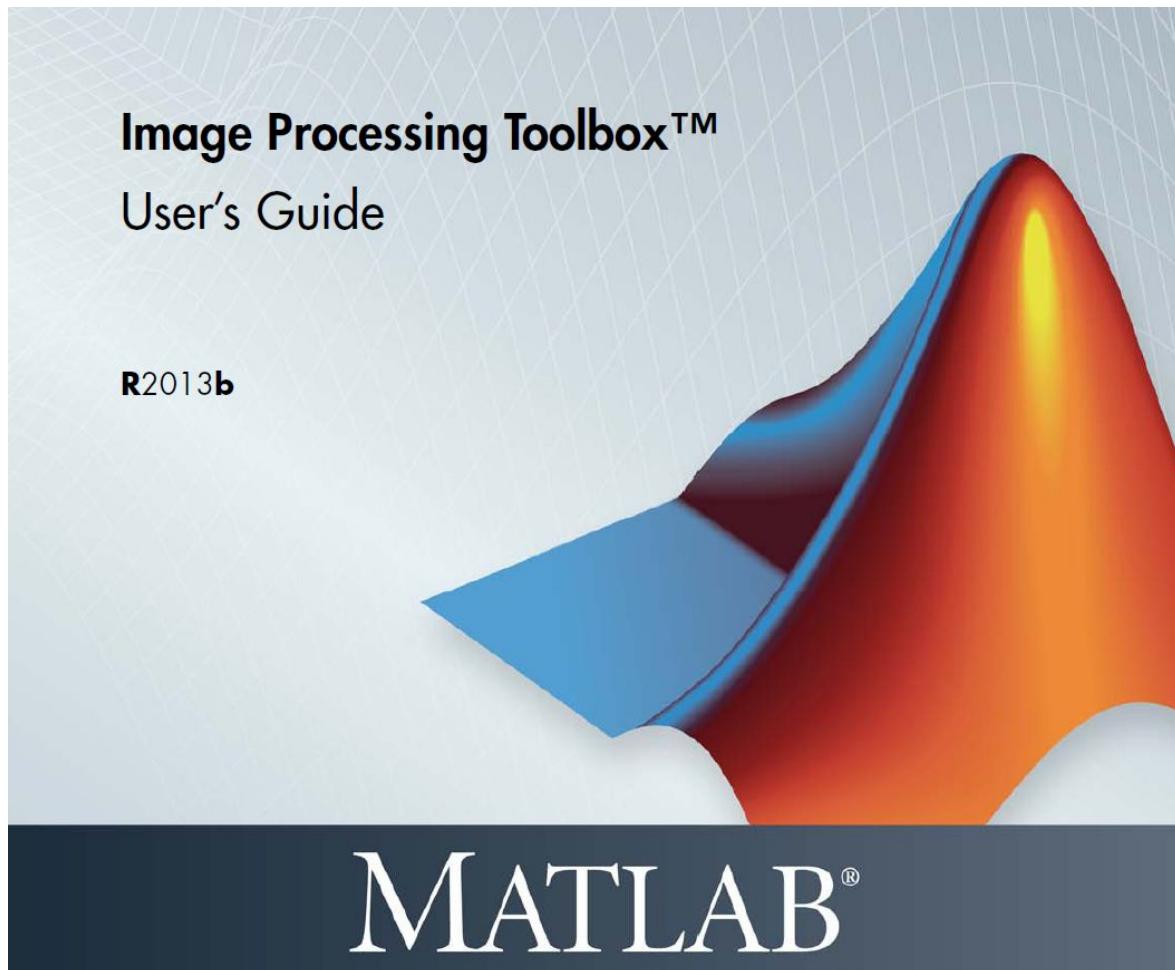


Gaussian white noise

# Basic Image Processing in Matlab



# Image Processing in Matlab (1)



# Image Processing in Matlab (2)

- The Matlab **Image Processing Toolbox** provides algorithms, functions, and apps for image processing and visualisation.
- Supports many image formats (JPG, TIF, PNG, BMP...)
- Tools for: image enhancement, filtering, analysis, segmentation, transformations, navigation, display...
- Check Moodle for MATLAB Download/Install info
- Check toolbox documentation!

# Image Processing in Matlab (3)

- Read an image:

```
I = imread('pout.tif');
```

- Display an image:

```
imshow(I);
```

- Save image file:

```
imwrite(I2, 'pout2.png');
```

- Convert RGB image to grayscale:

```
J = rgb2gray(I);
```

# Image Processing in Matlab (4)

- Image negative:

```
J = imcomplement(I);
```

- Image information:

```
imfinfo('pout2.png');
```

- Information about supported image formats:

```
imformats
```

- Resize image:

```
J = imresize(I,scale); //e.g. scale=1.5
```

# Image Processing in Matlab (5)

- Add noise:

```
J = imnoise(I, 'salt & pepper', 0.02);
```

- RGB to HSV/HSI models:

```
I_HSV = rgb2HSV(I_RGB);
```

- Rotate image:

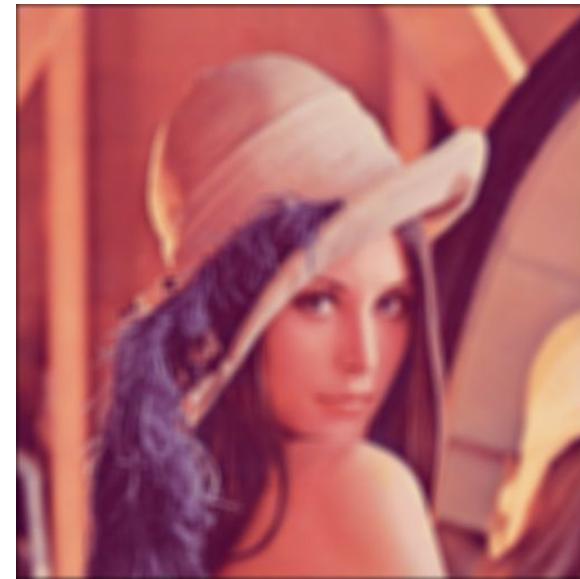
```
J = imrotate(I, angle);
```

- Crop image:

```
J = imcrop(I);
```



# Digital Image Filtering - Transforms

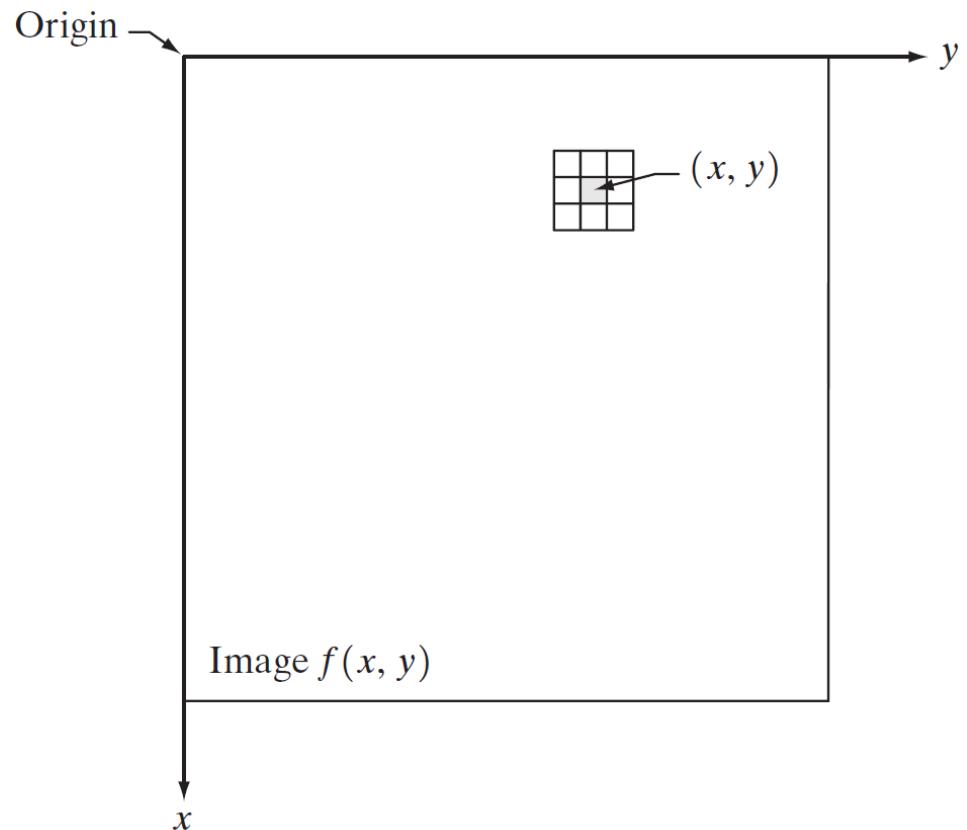


# Spatial Filtering



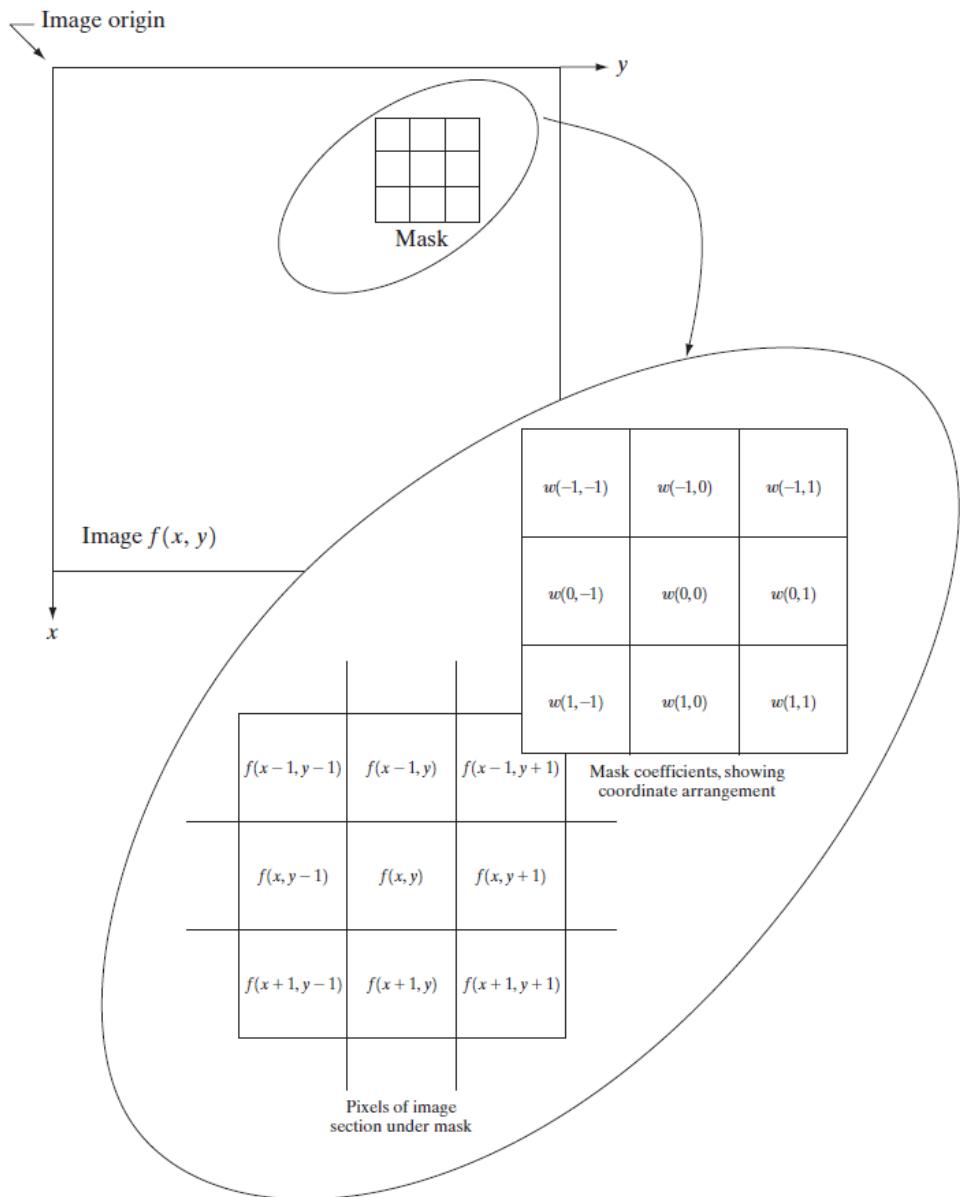
# Spatial Filtering (1)

- **Spatial domain:** pixels composing an image
- Spatial domain methods operate directly on pixels
- Spatial filtering typically uses neighborhood pixels
- **Figure:** Neighborhood about a point  $(x, y)$





- **Mask:** small sub-image used for spatial filtering
- Each mask has corresponding mask coefficients
- **Coefficients** determine the effect of the filtering process (e.g. sharpening, blurring)
- **Figure:** a 3x3 mask



# Spatial Filtering (3)

- Masks are sometimes called **convolution masks**
- Spatial filtering is directly related to **convolution**
- **Figure:** a 3x3 mask
- Mask output ( $z$ : pixel values):

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

$$R = \sum w_i z_i = w_1 z_1 + w_2 z_2 + \dots + w_9 z_9$$



# Smoothing Filters (1)

- Smoothing filters are used for **blurring** and **noise reduction**





# Smoothing Filters (2)

- **Smoothing linear filter**: its output is the average of the pixels contained in the neighborhood of the filter mask.
- Also called averaging filters or lowpass filters.
- The resulting image has reduced “sharp” transitions.
- This also results in blurry edges!



# Smoothing Filters (3)

- Figure: two 3x3 smoothing masks.

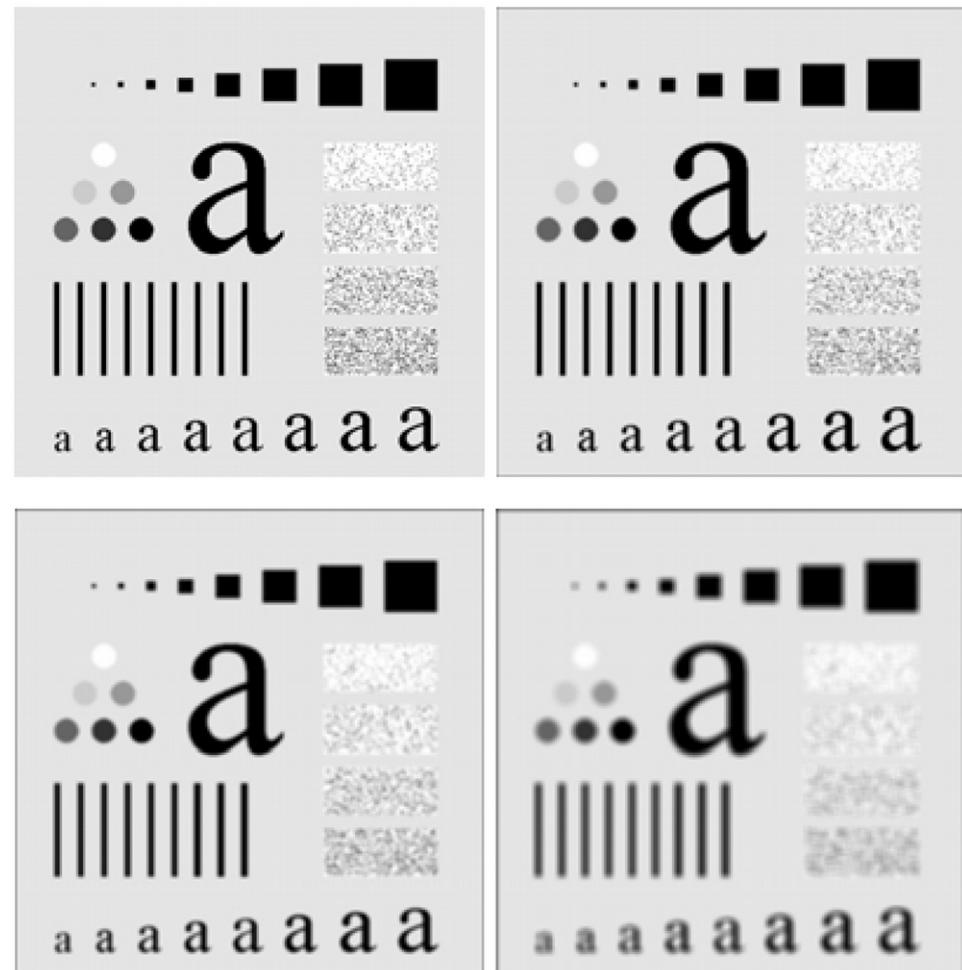
$\frac{1}{9} \times$	<table border="1"><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1	1	$\frac{1}{16} \times$	<table border="1"><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>2</td><td>4</td><td>2</td></tr><tr><td>1</td><td>2</td><td>1</td></tr></table>	1	2	1	2	4	2	1	2	1
1	1	1																			
1	1	1																			
1	1	1																			
1	2	1																			
2	4	2																			
1	2	1																			

- Left filter: average. Right filter: weighted average.



# Smoothing Filters (4)

- Figure: smoothing with averaging masks of sizes 3, 5, 9, and 11 pixels.
- Note the black border with high mask size: result of **zero padding**.
- **Zero padding:** expanding the image borders with 0's and then trimming off padded area.



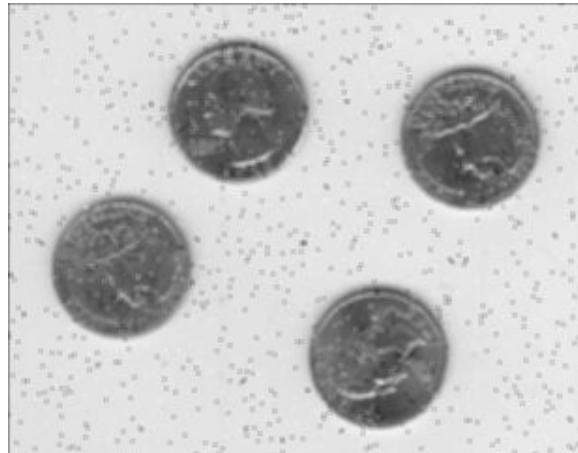
# Smoothing Filters (5)

- **Order-statistics filters**: filters whose response is based on ordering the pixels contained in the filter area.
- Most common order-statistics filter: **median filter**.
- Median filters have excellent noise reduction capabilities, with less blurring.
- **Example**: given a 3x3 neighborhood with values (10, 20, 20, 20, 15, 20, 20, 25, 100)  
By sorting these values, the median is the 5th largest value, i.e. 20.



# Smoothing Filters (6)

- **Figure:** Noisy image, average-filtered, median-filtered





# Sharpening Filters (1)

- **Objective:** to highlight fine detail or enhance detail that has been blurred.



# Sharpening Filters (2)

- Image sharpening can enhance edges and discontinuities (including noise).
- Since smoothing can be achieved by pixel summing, it is logical that sharpening can be achieved by **differentiation**.
- The **Laplacian** is the most common sharpening filter.
- The Laplacian is based on **second-order difference**:
  - First-order difference (1-D):  $f(x+1) - f(x)$
  - Second-order derivative (1-D):  $f(x+1) + f(x-1) - 2f(x)$
- First-order difference are suitable for edge enhancement
- Second-order derivatives have the ability to enhance fine detail.



# Sharpening Filters (3)

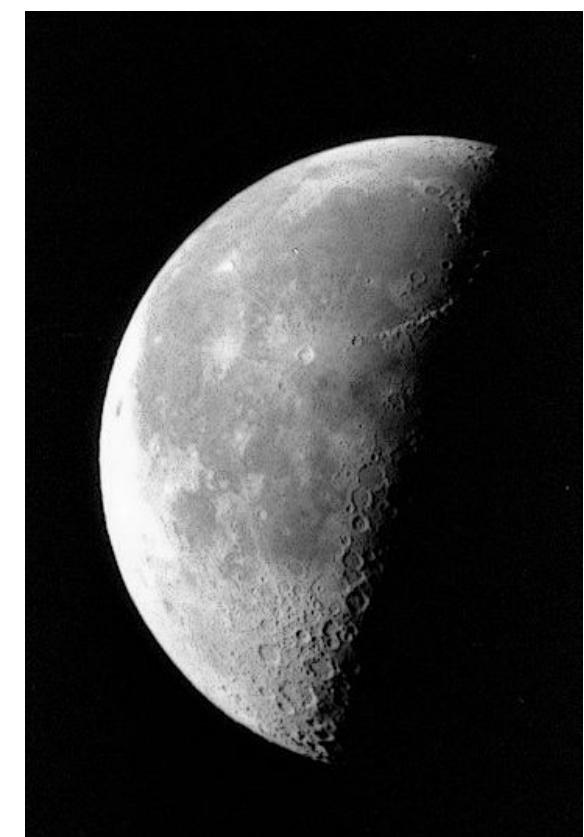
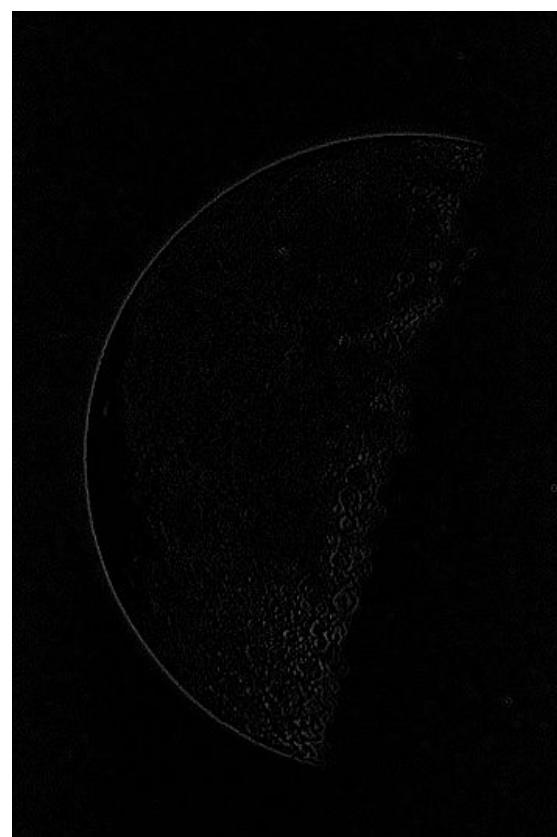
- The Laplacian highlights discontinuities and edges.
- **Figure:** different implementations of the Laplacian.

0	1	0	1	1	1
1	-4	1	1	-8	1
0	1	0	1	1	1
0	-1	0	-1	-1	-1
-1	4	-1	-1	8	-1
0	-1	0	-1	-1	-1



# Sharpening Filters (4)

- For detail enhancement, add the Laplacian from the original image (if the center coefficient is positive).



# Sharpening Filters (5)

- Instead of performing image enhancement in two steps (computing mask and then subtracting from image), this can be done in one step, with the following mask(s):

0	-1	0
-1	5	-1
0	-1	0

-1	-1	-1
-1	9	-1
-1	-1	-1

# Sharpening Filters (6)

- Other sharpening method: using first derivatives
- Sobel masks (for horizontal and vertical edges):

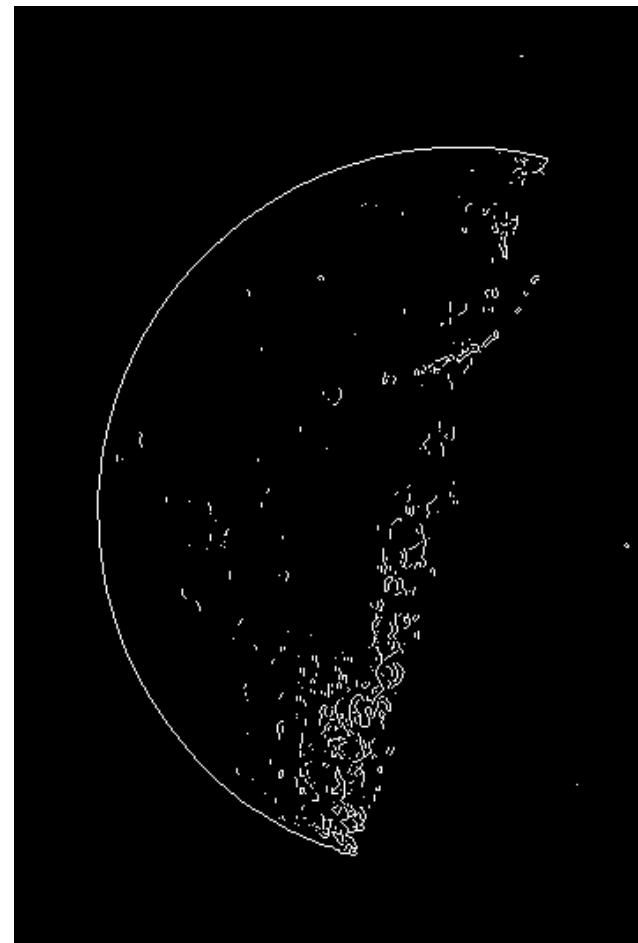
-1	-2	-1
0	0	0
1	2	1

-1	0	1
-2	0	2
-1	0	1



# Sharpening Filters (7)

- By combining output of vertical and horizontal filters, we can perform edge detection.
- **Figure:** edges computed using Sobel masks.



# 2-D Convolution (1)

- The process of moving a mask from pixel to pixel, and compute a quantity at each pixel, is the foundation of the **convolution** process.
- Given image  $f[x,y]$  and mask  $h[x,y]$ , the 2-D convolution is defined as:

$$f[x, y] * h[x, y] = \sum_m^{M-1} \sum_n^{N-1} f[m, n] h[x-m, y-n]$$

- When given a filter for convolution, we need to rotate the mask 180 degrees



# 2-D Convolution (2)

Example:

Image

2	2	2
2	2	2
2	2	2

Mask

1	2	1
2	1	2
1	2	1

Darker gray area in mask denotes its centre.

Output

2	6	8	6	2
6	12	18	12	6
8	18	26	18	8
6	12	18	12	6
2	6	8	6	2

Darker gray area represents output if we crop to original image size.

# 2-D Convolution (3)

- As in 1-D case, the 2-D convolution is **symmetric**:

$$f[x,y] * h[x,y] = h[x,y] * f[x,y]$$

- In image processing, we usually compute convolutions in the frequency domain instead of the spatial domain (next topic)

# Filtering in the Frequency Domain

# Filtering in the Frequency Domain

- Most image filtering is done in the frequency domain
- Main reason: computational efficiency.
- Also: frequency transforms (and especially the Fourier transform) have strong mathematical foundations and properties that make image processing faster & easier.

# 1-D Discrete Fourier Transform

- In previous weeks, we saw the one-dimensional **discrete Fourier transform (DFT)**:

$$X[k] = \sum x[n] e^{-i \frac{2\pi}{N} nk}$$

- $X$  is the **spectrum** of **signal**  $x$
- **Coefficients** of finite **combination** of **complex sinusoids**, ordered by frequencies
- Inverse Fourier transform returns original signal

# 2-D Discrete Fourier Transform (1)

- Extending the 1-D DFT into two dimensions is straightforward:

$$X[k, l] = \sum \sum x[m, n] e^{-i2\pi(km/M + ln/N)}$$

- Similarly, we can obtain the inverse 2-D DFT:

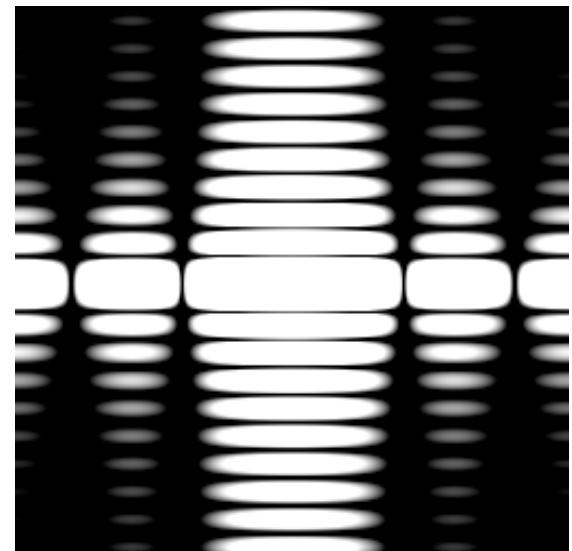
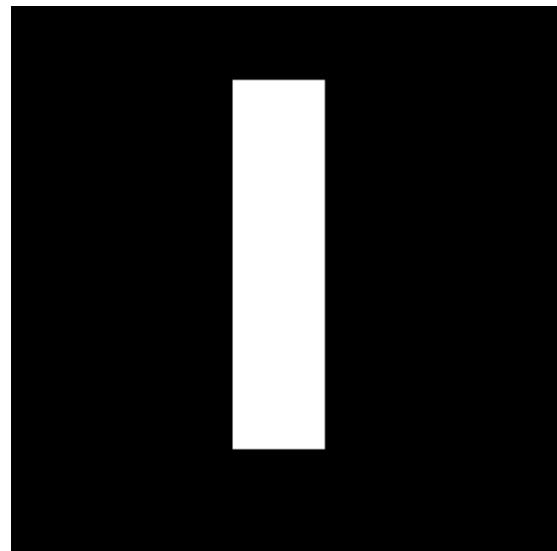
$$x[m, n] = \frac{1}{MN} \sum \sum X[k, l] e^{i2\pi(km/M + ln/N)}$$



# 2-D Discrete Fourier Transform (2)

- As in the 1-D case,  $X(k,l)$  is complex.
- $X(k,l)$  gives the magnitude (absolute) and phase (angle, arg) for each frequency in the continuum and is called the spectrum of  $x[m,n]$ .

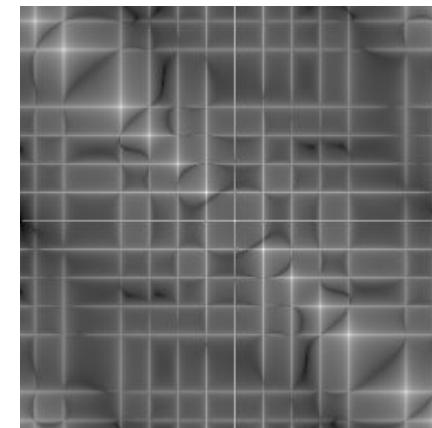
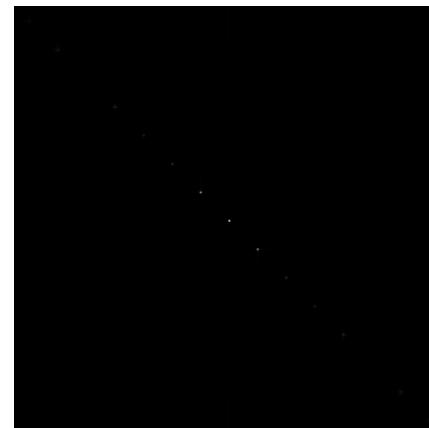
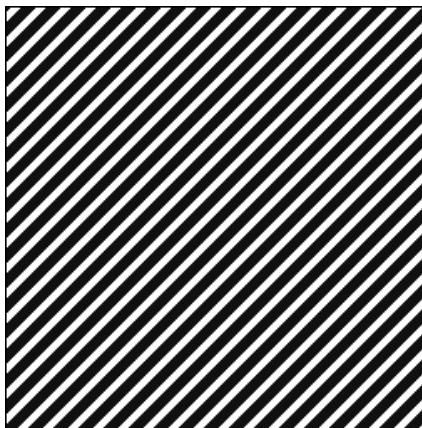
Figure: rectangle  
and its DFT





# 2-D Discrete Fourier Transform (3)

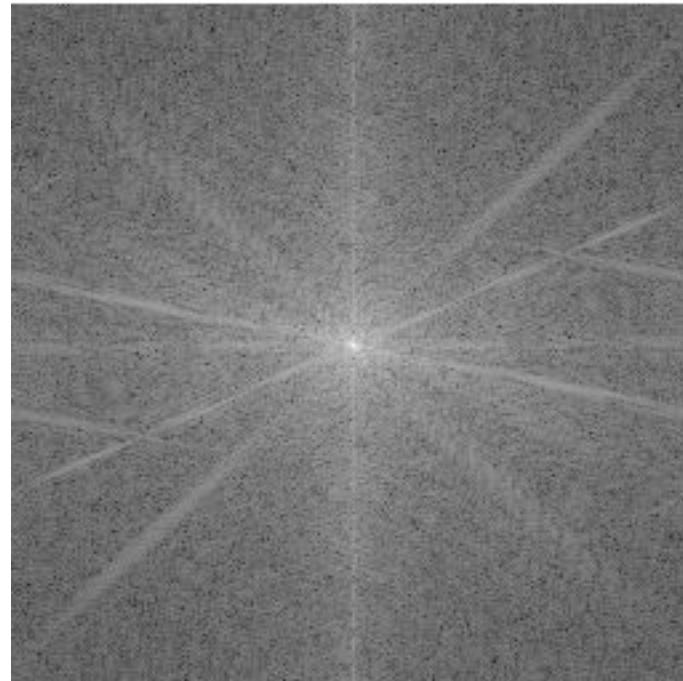
- It is not very easy to make direct associations between an image and its transform...
- But: we can make some general statements about the frequency components of an image. **Figure**: image, DFT, and logarithm of the DFT.





# 2-D Discrete Fourier Transform (4)

- A more realistic example:



# 2-D Discrete Fourier Transform (5)

- **Convolution Theorem:** convolution in the spatial domain can be expressed by multiplication in the frequency domain:

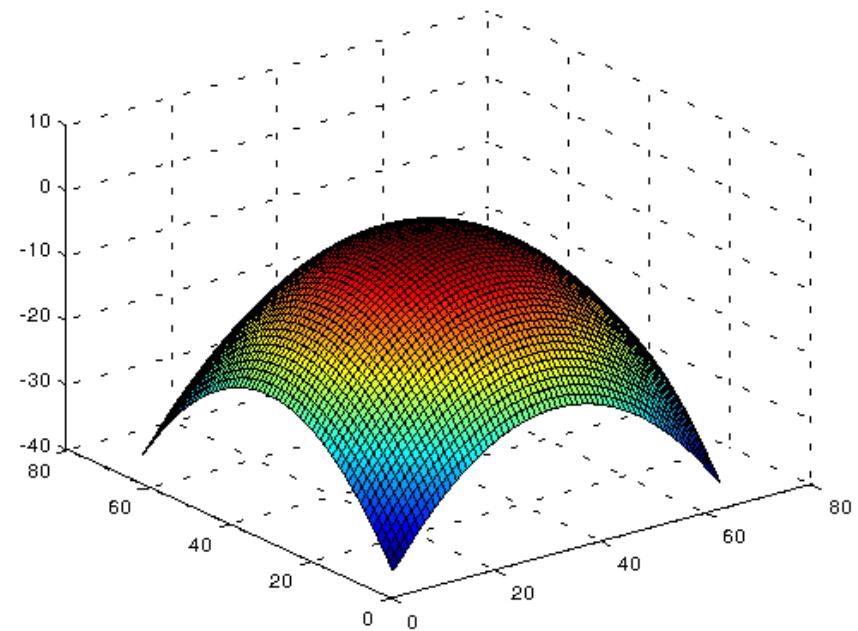
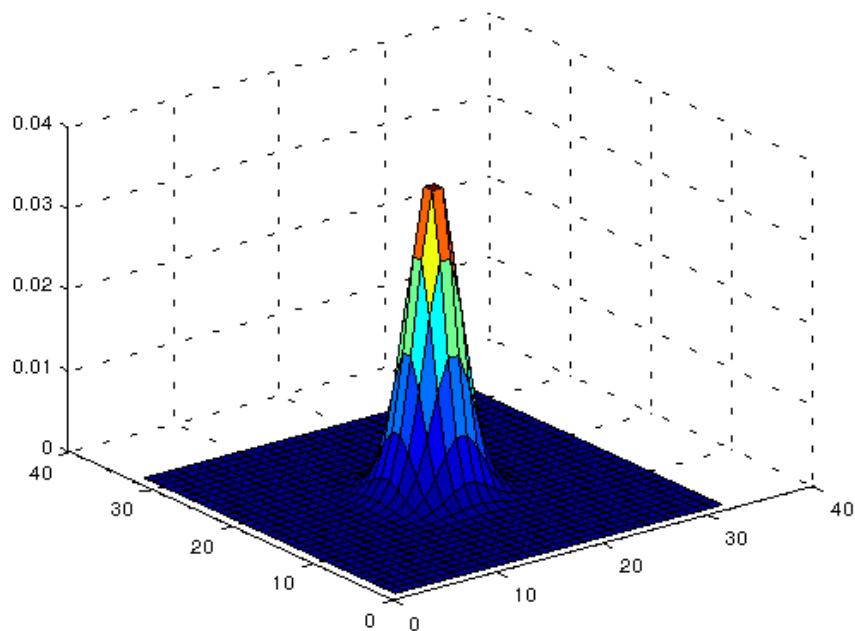
$$f[x,y] * h[x,y] \Leftrightarrow F[k,l] H[k,l]$$

- In practice when filtering, we compute the DFT of an image and its filter. We multiply them, and compute the inverse DFT.
- **Fast Fourier Transform (FFT):** fast implementation.  $O(M \log M)$  instead of  $O(M^2)$ .



# Frequency Domain Filters (1)

- **Figure:** Gaussian mask and its DFT (lowpass filter)

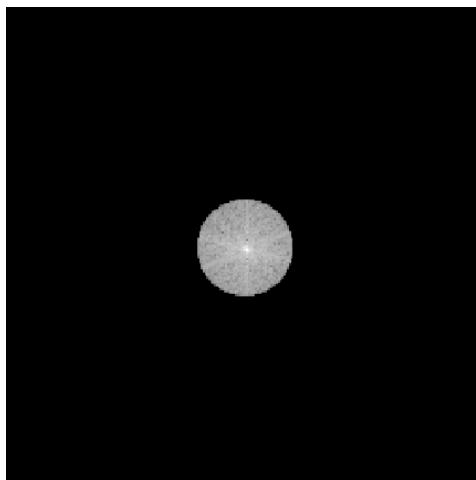
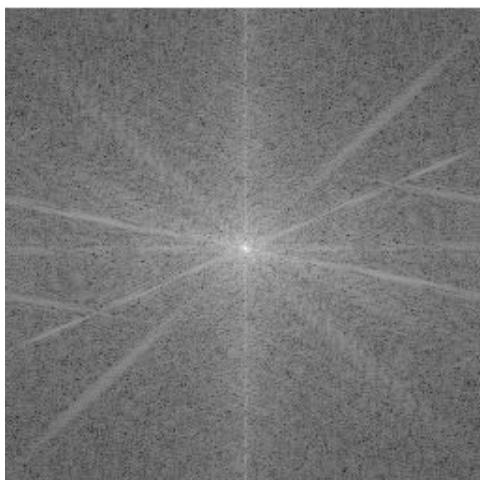




# Frequency Domain Filters (2)



**Figure:** lowpass  
filtering example.

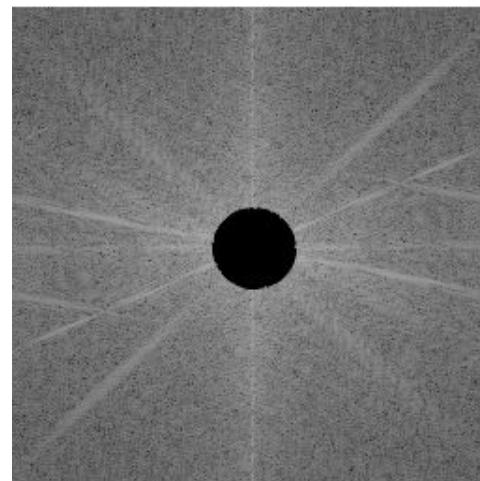
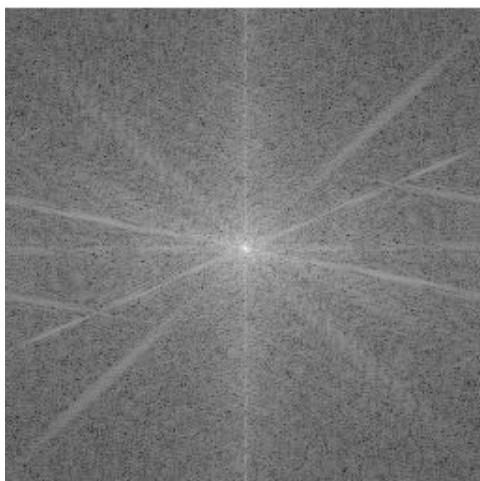




# Frequency Domain Filters (3)



**Figure:** highpass  
filtering example.



# Frequency Domain Filters (4)

- **Smoothing frequency-domain filters** (ordered from sharp to smooth cut-offs): ideal, Gaussian, Butterworth.
- **Sharpening frequency-domain filters**: same types as smoothing filters, typically their complement.



# Image Filtering and Transforms in Matlab

# Filtering and Transforms in Matlab (1)

- Define averaging filter:

```
h = ones(5,5) / 25;
```

- Filter an image:

```
J = imfilter(I,h);
```

- Use pre-existing filter (Gaussian, Laplacian, Sobel...):

```
h = fspecial(type);
```

- Compute frequency response for filter:

```
H = freqz2(h,n1,n2); //e.g. n1,n2=256
```



# Filtering and Transforms in Matlab (2)

- Perform median filtering:

```
J = medfilt2(I, [3 3]);
```

- 2-D FFT:

```
F = fft2(I);
```

- We typically visualise the logarithm of the magnitude:

```
F2 = log(abs(F)); //try imagesc for plotting
```

- Shift zero-frequency component to center of spectrum:

```
F2 = fftshift(F);
```

# Filtering and Transforms in Matlab (3)

- Edge detection:

```
BW = edge(I, 'sobel');
```

- Inverse 2-D FFT:

```
B = ifft2(F);
```

- 2-D convolution:

```
C = conv2(A, B);
```

# Properties of the 2-D Fourier Transform

# 2-D DFT Properties (1)

- Similar properties with the 1-D Fourier Transform

- **Linearity:**

$$a x_1 + b x_2 \Leftrightarrow a X_1 + b X_2$$

- **Scaling:**

$$x[ak, bl] \Leftrightarrow \frac{1}{|ab|} X[k, l]$$

## 2-D DFT Properties (2)

- Separability:

$$X[k, l] = \frac{1}{M} \sum_{m=0}^{M-1} X[m, l] e^{-j2\pi lm/M}$$

- Where:

$$X[m, l] = \frac{1}{N} \sum_{n=0}^{N-1} x[m, n] e^{-j2\pi ln/N}$$

- In other words: we can compute the 2-D transform by first computing a 1-D transform over each row, and then computing a 1-D transform over each column.



# 2-D Correlation - Template Matching

# 2-D Correlation (1)

- As in the 1-D case, the **correlation** measures the similarity of two signals by adding their products at each sampled time.
- **Definition of 2-D correlation:**

$$x[m, n] \circ y[m, n] = \frac{1}{RS} \sum_{r=0}^{R-1} \sum_{s=0}^{S-1} x^*[r, s] y[m+r, n+s]$$

where  $*$  denotes complex conjugate (since we deal with real signals tough,  $x^* = x$ )

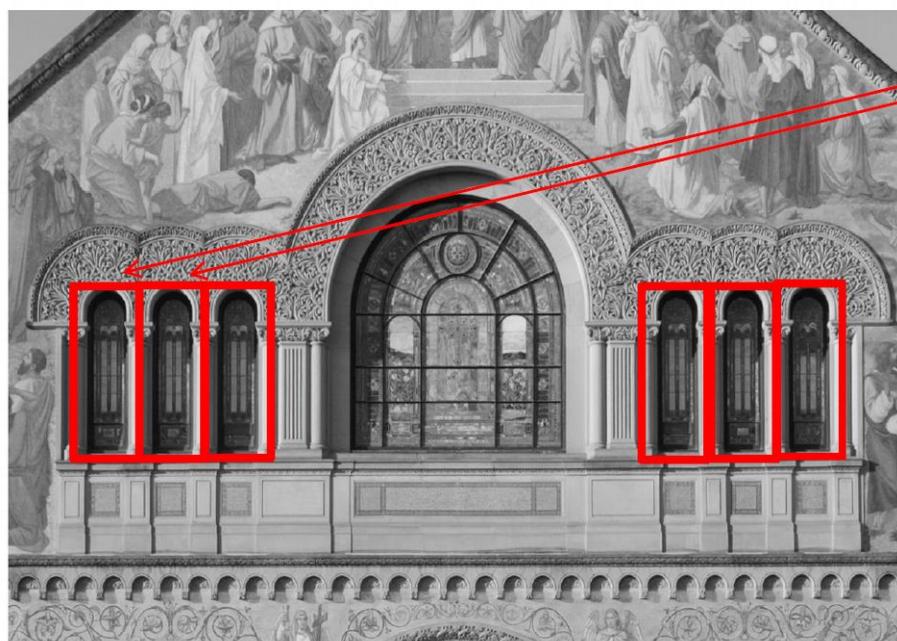
## 2-D Correlation (2)

- The correlation is very similar to the convolution (with the exception of the conjugate and the positive sign)
- It can be implemented using convolution in Matlab:  
`conv2(a, rot90(conj(b), 2))`
- The positive sign means that we don't mirror the mask about the origin.
- **Correlation theorem:**  
$$x[m, n] \circ y[m, n] \Leftrightarrow X^*[k, l] Y[k, l]$$



# Template Matching (1)

- The principal use of correlation is for matching (i.e. locate an object)



$t[x,y]$

$s[x,y]$

# Template Matching (2)

- Use of correlation for template matching:

The term watershed  
refers to a ridge that ...

... divides areas  
.. drained by different  
river systems.

a

- 2-D Correlation –  
Template Matching
- Directly done in Matlab  
with 2d cross-correlation:

`C = xcorr2 (A, B);`



# Template Matching (3)

- Left: 2-D correlation. Right: thresholded correlation.

