

Computer Vision

INM460/IN3060

Lecture 3

Image filtering and edge detection

Dr Giacomo Tarroni

Slides credits: Giacomo Tarroni, Sepehr Jalali, Wenjia Bai

Admin

- Drop-in hours for Week 4: Wednesday 12/02, 11am – 1pm, A302B
- Others remain as default: Wednesdays 2 – 4pm, A302B

Recap from the previous lecture

- Digital images:
 - Formation
 - Pinhole camera model
 - Digital representation
- Digital image processing:
 - Brightness and colour transformations
 - Geometric transformations
 - Filtering (not yet presented)

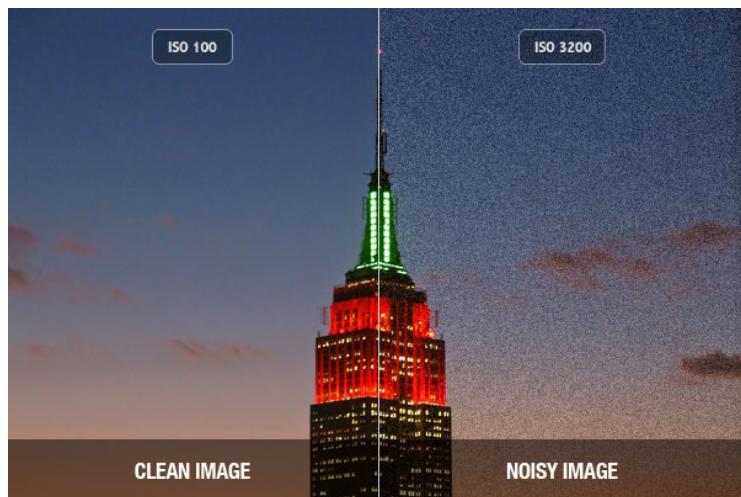
Overview of today's lecture

- Image filtering
- Linear filtering:
 - Convolution vs cross-correlation
 - Common filters:
 - Moving average
 - Gaussian
 - Sharpening
- Non-linear filtering: median filter
- Edge detection
 - Gradient-based
 - Laplacian-based
 - Non-maximum suppression
- Common CV tasks

Imperfections in images



Low
resolution



Noise



Bloom
(i.e. light bleeding on darker background)

Imperfections in images



Motion blur



Poor contrast



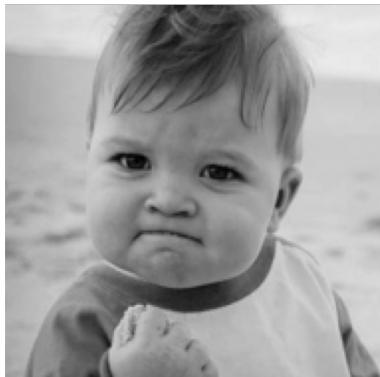
Compression artefacts



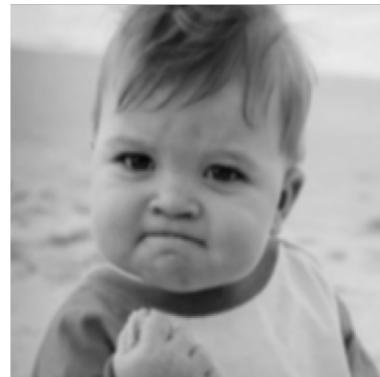
Lens distortion

Image filtering

- The easiest way to improve image appearance is to perform image filtering
- In general, filtering allows to remove unwanted components from a given signal or to enhance some desired ones
- Image filtering can be used in a variety of contexts, including noise reduction and image enhancement
- Filtering can be applied to both grayscale and colour images. For colour images, the filtering is applied to each channel separately or to the V channel after a conversion to HSV colour space



Original image



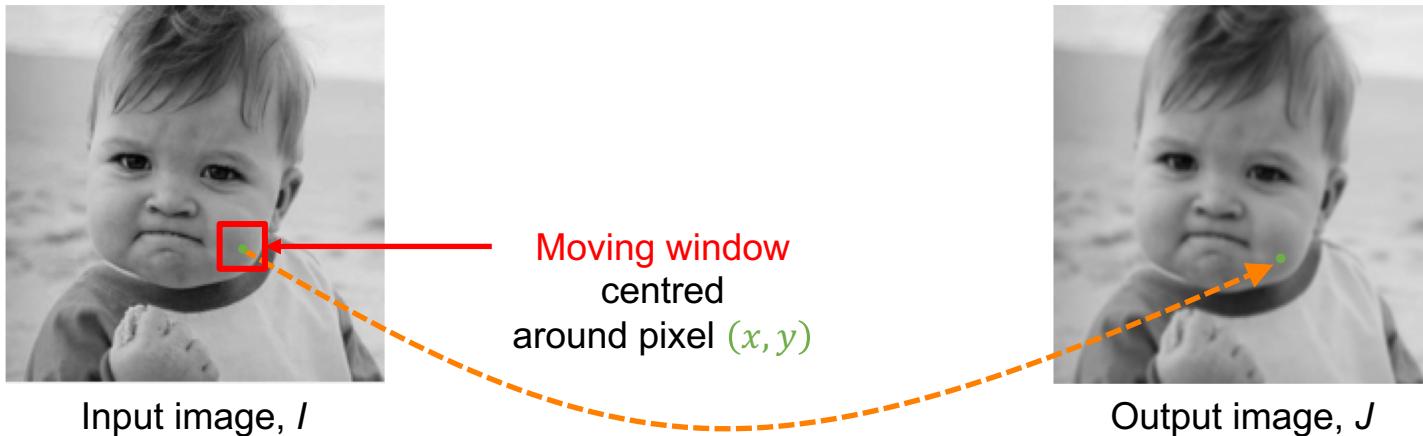
Smoothed
(-noise, -details)



Sharpened
(+details, +noise)

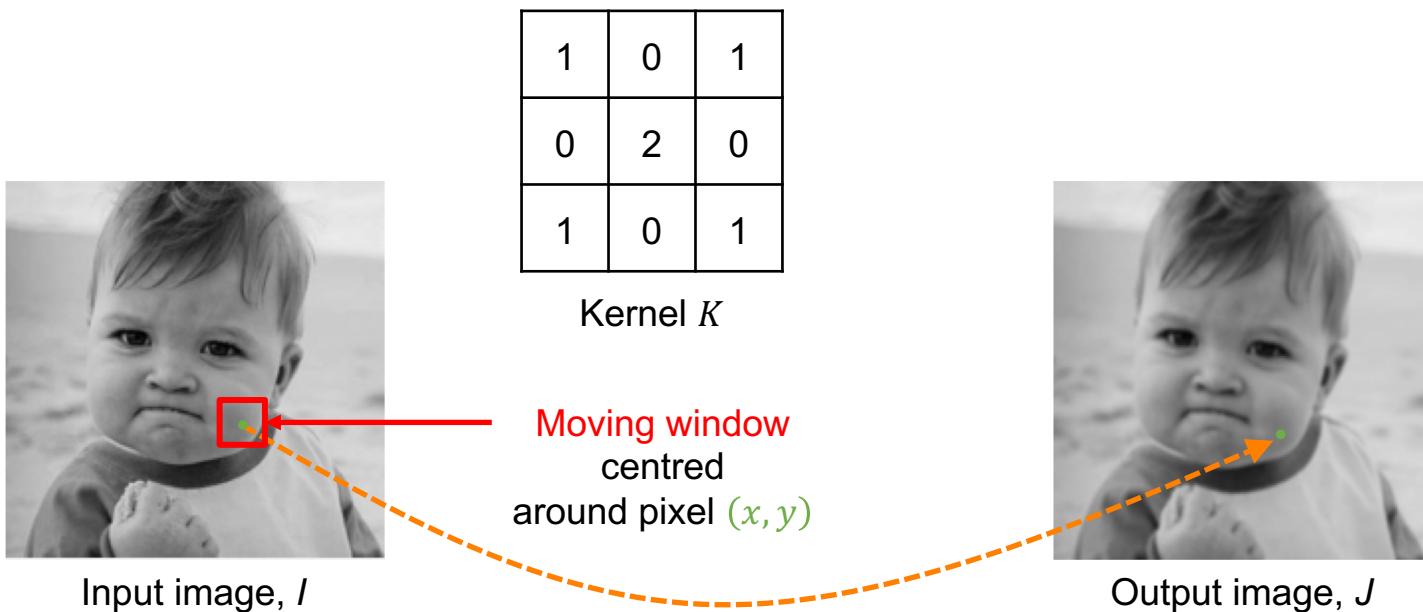
Window-based image filtering

- In window-based image filtering, the value of each pixel (x, y) in the output image J is a function of the pixel values of the input image I in a neighbourhood of (x, y)
- The neighbourhood is defined by a window of a given size (e.g. $N \times N$)
- This can be thought as sliding a $N \times N$ window on the input image, computing an output value at each position
- The type of filtering is determined by the function describing the relationship between input and output pixel values



Linear filtering

- If the pixel values in the output image are a **linear combination** of pixel values in a neighbourhood of the input image, then all we need to represent the function between the input and the output are the weights
- The weights are usually represented as a **kernel**
- If, in addition, the weights are fixed for all positions, we are performing **linear filtering**



Linear filtering

- Our goal is to compute $J(x, y)$, the value of pixel (x, y) in the output image J
- To do this, we place the kernel K (a matrix of weights) on the input image I centred around a position (x, y) . For every position (n, m) in K , we multiply the kernel value with the input image pixel, and sum up all the values:

$$J(x, y) = \sum_{n,m \in K} K(n, m) I(x + n, y + m)$$

- For example, consider the following 3x3 kernel, and a set of pixel values:

1	0	1
0	2	0
1	0	1

Kernel K

*

22	23	24
27	28	29
32	33	34

Input image I

=

?		

Output image J

Linear filtering

- The kernel is typically much smaller than the input image. The output image is generated by **sliding** the kernel on the input image, one pixel at a time

$$\begin{array}{|c|c|c|} \hline 1 & 0 & 1 \\ \hline 0 & 2 & 0 \\ \hline 1 & 0 & 1 \\ \hline \end{array} \quad * \quad
 \begin{array}{|c|c|c|c|c|} \hline 22 & 23 & 24 & 25 & 26 \\ \hline 27 & 28 & 29 & 30 & 31 \\ \hline 32 & 33 & 34 & 35 & 36 \\ \hline 37 & 38 & 39 & 40 & 41 \\ \hline 42 & 43 & 44 & 45 & 46 \\ \hline \end{array} \quad = \quad
 \begin{array}{|c|c|c|c|c|} \hline X & X & X & X & X \\ \hline X & 168 & & & X \\ \hline X & & & & X \\ \hline X & & & & X \\ \hline X & X & X & X & X \\ \hline \end{array}$$

Kernel K

Input image I

Output image J

Linear filtering

- The kernel is typically much smaller than the input image. The output image is generated by **sliding** the kernel on the input image, one pixel at a time

$$\begin{array}{|c|c|c|} \hline 1 & 0 & 1 \\ \hline 0 & 2 & 0 \\ \hline 1 & 0 & 1 \\ \hline \end{array}
 \quad *
 \quad
 \begin{array}{|c|c|c|c|c|} \hline 22 & 23 & 24 & 25 & 26 \\ \hline 27 & 28 & 29 & 30 & 31 \\ \hline 32 & 33 & 34 & 35 & 36 \\ \hline 37 & 38 & 39 & 40 & 41 \\ \hline 42 & 43 & 44 & 45 & 46 \\ \hline \end{array}
 \quad =
 \quad
 \begin{array}{|c|c|c|c|c|} \hline X & X & X & X & X \\ \hline X & 168 & 174 & & X \\ \hline X & & & & X \\ \hline X & & & & X \\ \hline X & X & X & X & X \\ \hline \end{array}$$

Kernel K

Input image I

Output image J

Linear filtering

- The kernel is typically much smaller than the input image. The output image is generated by **sliding** the kernel on the input image, one pixel at a time

$$\begin{array}{|c|c|c|} \hline 1 & 0 & 1 \\ \hline 0 & 2 & 0 \\ \hline 1 & 0 & 1 \\ \hline \end{array}
 \quad *
 \quad
 \begin{array}{|c|c|c|c|c|} \hline 22 & 23 & 24 & 25 & 26 \\ \hline 27 & 28 & 29 & 30 & 31 \\ \hline 32 & 33 & 34 & 35 & 36 \\ \hline 37 & 38 & 39 & 40 & 41 \\ \hline 42 & 43 & 44 & 45 & 46 \\ \hline \end{array}
 \quad =
 \quad
 \begin{array}{|c|c|c|c|c|} \hline X & X & X & X & X \\ \hline X & 168 & 174 & 180 & X \\ \hline X & 198 & 204 & 210 & X \\ \hline X & 228 & 234 & 240 & X \\ \hline X & X & X & X & X \\ \hline \end{array}$$

Kernel K Input image I Output image J

Linear filtering

- Kernel normalization:
Each kernel element is **divided by the sum of all kernel elements (with their sign!)**, so that the sum of the elements of a normalized kernel is unity. This ensures that the intensity average for the whole image remains unchanged through filtering
- Clipping due to datatype ranges:
If datatype is uint8, then range of values is 0÷255, which might cause unwanted clipping. When using filters, first convert your image to double, then do the calculations. Before converting it back to uint8 normalize the values

Cross-correlation vs convolution

- The act of moving a linear kernel on the input image to create the output one is called cross-correlation. Another very similar operation is convolution
- Cross-correlation:

$$J(x, y) = K \star I = \sum_{n,m} K(n, m) I(x + n, y + m)$$

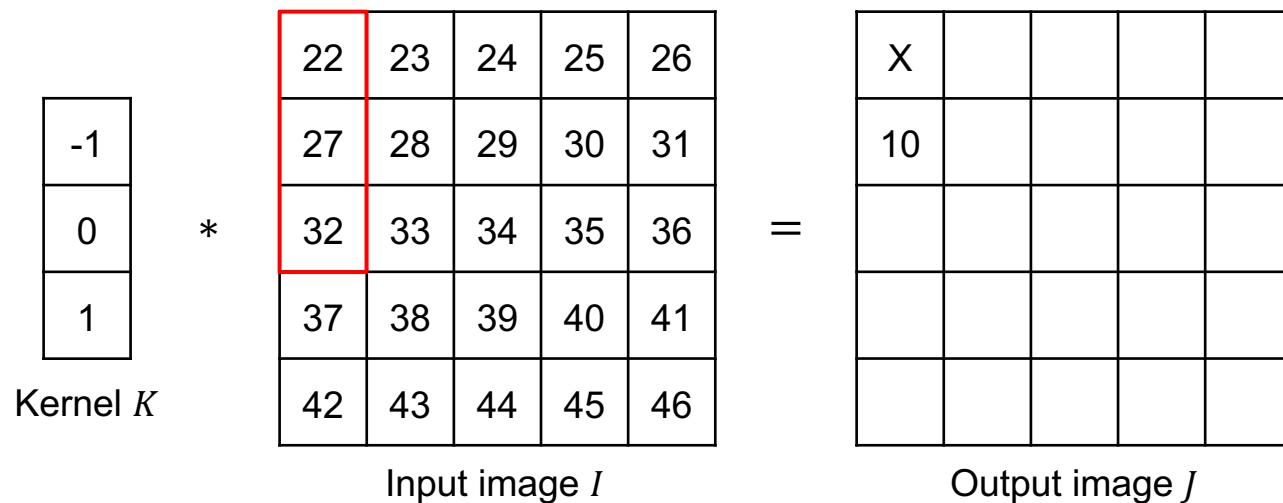
- Convolution:
- $$J(x, y) = K * I = \sum_{n,m} K(n, m) I(x - n, y - m) = \sum_{n,m} K(-n, -m) I(x + n, y + m)$$
- Convolution is cross-correlation with a kernel flipped in both directions:
www.youtube.com/watch?v=Ma0YONjMZLI
 - They are the same for symmetric kernels
 - In practice, the term used is “convolution” but the operation we will actually perform is cross-correlation

Convolution

Convolution:

$$J(x, y) = K * I = \sum_{n,m} K(n, m) I(x - n, y - m)$$

- Convolution is not restricted to 2D kernels:



Convolution

Convolution:

$$J(x, y) = K * I = \sum_{n,m} K(n, m) I(x - n, y - m)$$

- Convolution is a linear operator. Thus, these properties hold:

- Commutative: $K * I = I * K$
- Distributive: $(K_1 + K_2) * I = K_1 * I + K_2 * I$
- Associative: $K_2 * (K_1 * I) = (K_2 * K_1) * I$

The associative property is useful to decompose 2D kernels into 2, 1D ones:

$$\begin{array}{|c|} \hline 1/3 \\ \hline 1/3 \\ \hline 1/3 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1/3 & 1/3 & 1/3 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline \end{array}$$

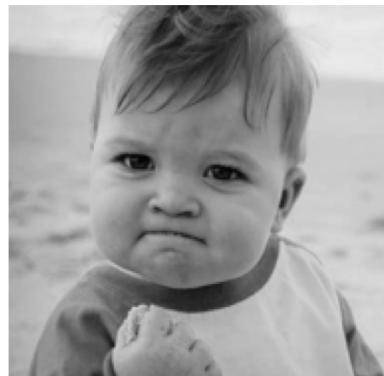
Identity

- The kernel of an identity filter is simply a set of zeros with a 1 in the centre
- The output pixel value will be equal to the input pixel value

0	0	0
0	1	0
0	0	0

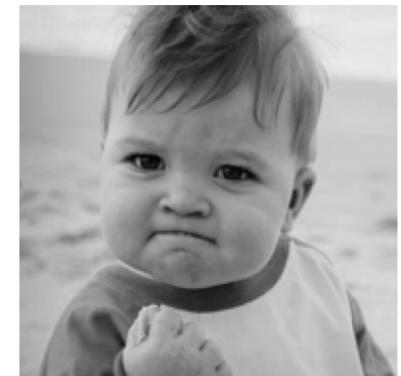
Kernel K
Identity

*



Input image I

=



Output image J

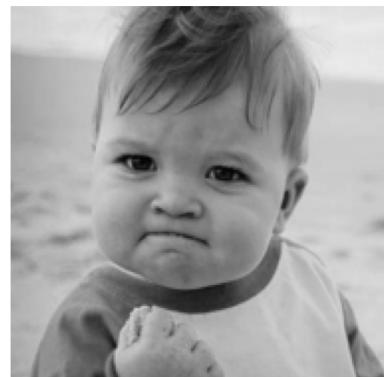
Moving average

- The kernel of a moving average (or box) filter is simply a set of ones divided by the number of kernel elements
- The output pixel value will be the **average** of the pixel values in the box
- Used for smoothing and blurring (noise reduction)

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

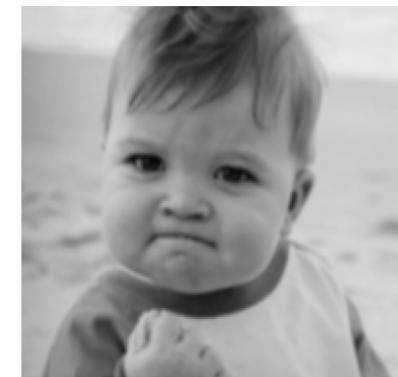
Kernel K
Moving average

*



Input image I

=



Output image J

Moving average

To implement this in Matlab, we must:

- Load the image I
- Define the kernel K
- Apply the filter to the image using the built-in function `imfilter`

```
% Load, resize, and show original image
I = double(imread('Success-Kid.jpg'));
I = imresize(I, 0.33);
figure; imshow(I/255); title('Original image');

% Create the kernel
K = (1/9) * ones(3,3);

% Apply to each colour channel and display result
J = imfilter(I, K);
figure; imshow(J/255); title('Filtered image');
```

- By default `imfilter` performs zero padding
- Other edge handling options are available



Moving average

What is the impact of the kernel size on the input?



$K = (1/9) * \text{ones}(3,3);$



$K = (1/25) * \text{ones}(5,5);$



$K = (1/49) * \text{ones}(7,7);$

Sharpening filter

Using the distributive property of convolution, we can see that the sharpening filter adds “image details” (i.e. differences from local average)

$$\begin{array}{|c|c|c|} \hline -1/9 & -1/9 & -1/9 \\ \hline -1/9 & 17/9 & -1/9 \\ \hline -1/9 & -1/9 & -1/9 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \begin{array}{|c|c|c|} \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline \end{array}$$

Kernel K
 Sharpening Input image Input image Smoothed image

“Image details”



Non-linear filtering

- So far, we have described linear filtering operations, which are mathematically represented by the convolution of a kernel with the image
- A different class of filtering methods are instead non-linear, and cannot be represented through convolution
- An example is the median filter

Median filter

Built-in Matlab function: `medfilt2`



Original image
(with added “salt & pepper” noise)



Median filtered image

Hand-crafted vs learned filters

- The filters seen so far are hand-crafted: their design (including the choice of parameters like the size of the kernel) is usually based on heuristic
- Machine learning allows to “learn” the filter parameters from a set of input and output images (e.g. noisy images as input and smoothed images as output)
- This aspect is particularly evident in deep learning techniques

Edge detection

- Edge detection is a fundamental problem in CV. It consists in locating the edges, i.e. strong discontinuities in pixel intensity and/or colour in the image
- This can be useful for distinguishing objects from their surroundings (both for humans and for computers)

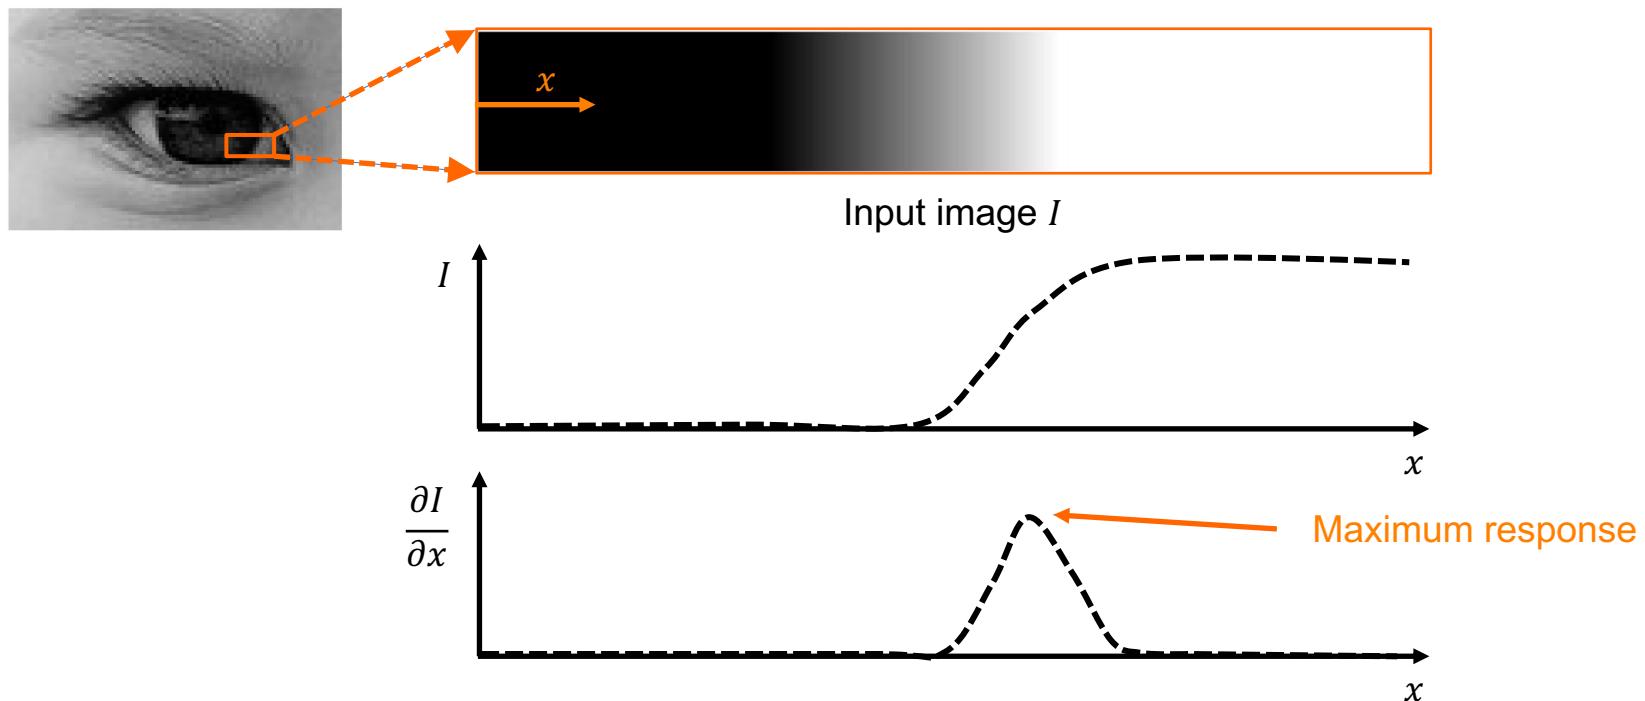


Edge detection

- A naïve edge detection algorithm is as follows. For each pixel in the image:
 - Analyse each of the 8 pixels surrounding it
 - Detect both the darkest and the lightest pixels
 - If $(I_{lightest} - I_{darkest}) > threshold$, then set the output pixel to 1; else set it to 0
- There are many techniques to perform it, including
 - gradient-based
 - Laplacian-based
- Usually performed on grayscale images

Gradient-based edge detection

- Edge detection can be performed using first spatial derivatives of the image intensity, i.e. evaluating the **(absolute) slope of image intensity** at each pixel (and for each dimension, x and y)
- High value of the first derivative along x for one pixel
 - high change in intensity along x
 - high chances of having an edge along x in that pixel



First derivative

- The first derivative can be approximated using a Taylor expansion:

$$I(x + h, y) = I(x, y) + h \frac{\partial I(x, y)}{\partial x} + O(h^2) \quad \text{Forward scheme difference}$$

$$I(x - h, y) = I(x, y) - h \frac{\partial I(x, y)}{\partial x} + O(h^2) \quad \text{Backward scheme difference}$$

- If we subtract the two, and consider the limit for $h \rightarrow 0$, we can write

$$I(x + h, y) - I(x - h, y) = 2h \frac{\partial I(x, y)}{\partial x}$$

$$I_x \equiv \frac{\partial I(x, y)}{\partial x} = \frac{I(x + h, y) - I(x - h, y)}{2h} \quad \text{Central scheme difference}$$

- The smallest h we can take is one pixel: to implement the derivative of $I(x, y)$ with respect to x at a pixel (x, y) , we simply take the intensity of the pixel to the right, subtract the intensity of the pixel to the left, and divide by 2
- For y , we have

$$I_y \equiv \frac{\partial I(x, y)}{\partial y} = \frac{I(x, y + h) - I(x, y - h)}{2h}$$

First derivative kernels

- A simple first derivative kernel with respect to x is

$$0.5 \times \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array}$$

First derivative kernel
(x direction)

```
K = 0.5*[-1 0 1];  
Ix = imfilter(I,K);  
imshow(Ix,[ ]);
```



I_x

First derivative kernels

- A simple first derivative kernel with respect to y is

$$0.5 \times \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

First derivative kernel
(y direction)

Transpose

```
K = 0.5*[-1 0 1]';  
Iy = imfilter(I,K);  
imshow(Iy,[ ]);
```



I_y

Image gradient

- The image gradient ∇I is a vector consisting of the two first derivatives computed at each point in the image

$$\nabla I(x, y) \equiv (I_x, I_y) = \left(\frac{\partial I(x, y)}{\partial x}, \frac{\partial I(x, y)}{\partial y} \right)$$

- Built-in Matlab function: `imgradient`

- It is usually represented with two images:

- Magnitude:

$$|\nabla I(x, y)| = \sqrt{I_x^2 + I_y^2}$$

$$|\nabla I(x, y)|$$

- Direction/angle:

$$\theta(x, y) = \arctan(I_y, I_x)$$

$$\theta(x, y)$$



Common first derivative kernels

- Derivatives are highly sensitive to noise in the images
- It is common to first smooth the image and then apply the derivative kernel:
 - Prewitt filter:

$$\begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

First derivative kernel
(x direction)

Smoothing kernel

Prewitt kernel
(x direction)

- Sobel filter:

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

First derivative kernel
(x direction)

Smoothing kernel

Sobel kernel
(x direction)

Derivative of Gaussian

- The Gaussian filter is also a common choice for smoothing
- Using the associative property of convolution

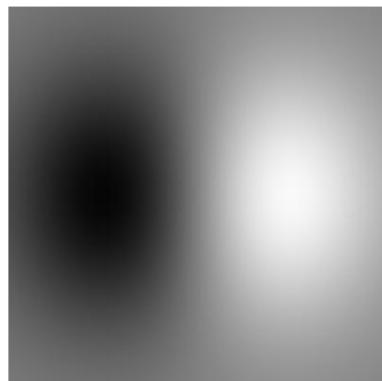
$$K_2 * (K_1 * I) = (K_2 * K_1) * I$$

we can create a single kernel that performs two operations:

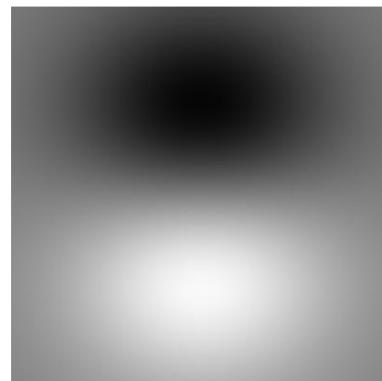
$$K_1 = \frac{1}{2\pi\sigma^2} e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$

$$K_2 = \frac{\partial}{\partial x}$$

- $K_3 = K_2 * K_1$ is the Derivative of Gaussian kernel (along x)



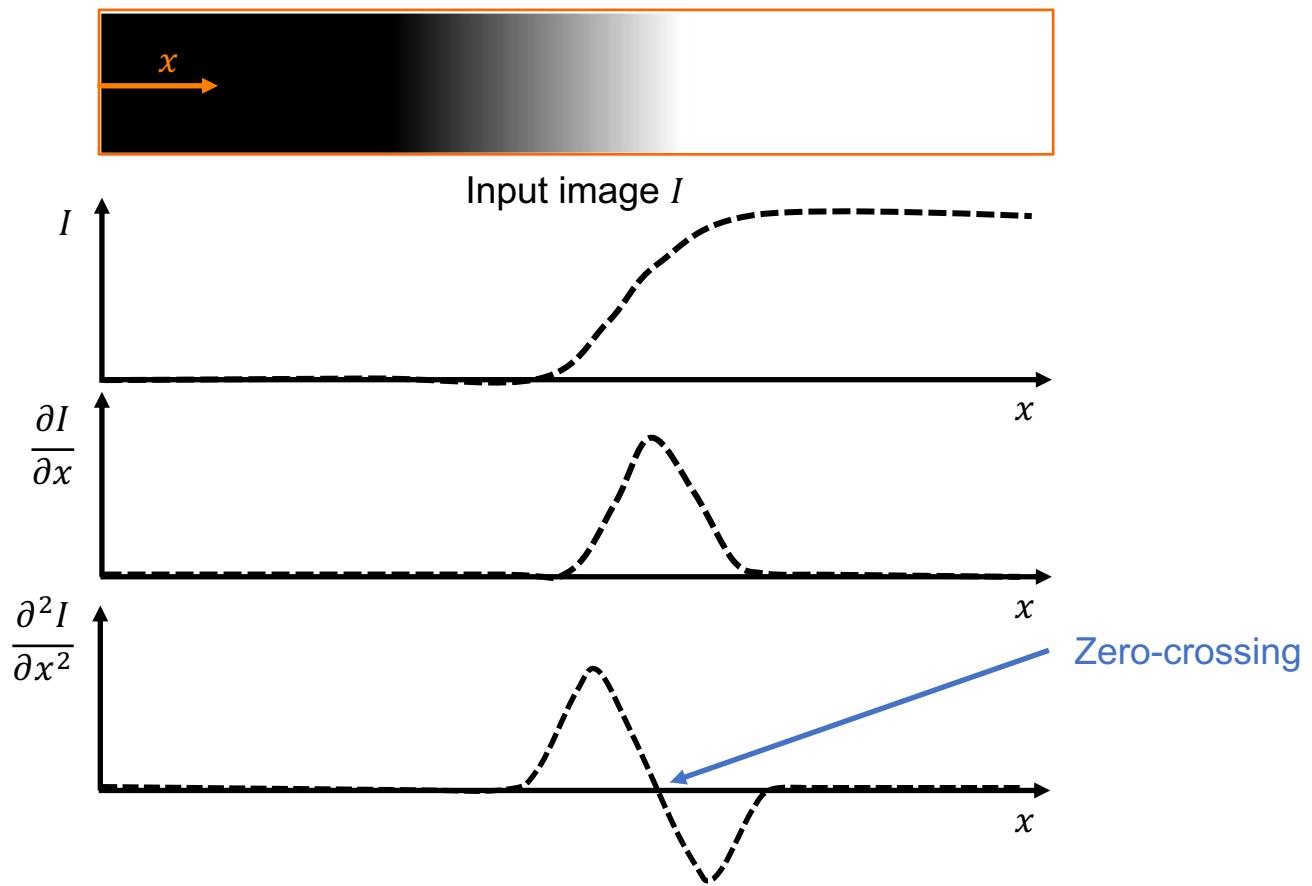
Derivative of Gaussian
(x direction)



Derivative of Gaussian
(y direction)

Laplacian-based edge detection

Edge detection can be performed using second spatial derivatives of the image intensity, i.e. looking for their **zero-crossings** at each pixel (along both dimensions)



Second derivative kernels

- If we consider the limit for $h \rightarrow 0$, the second derivative can be approximated with the following:

$$\frac{\partial^2 I(x, y)}{\partial x^2} = \frac{I(x + h, y) - 2I(x, y) + I(x - h, y)}{h^2}$$

- Since we are looking for the zero-crossings, we can sum the two derivatives
- The sum of the two second derivatives is the **Laplacian** ΔI of the image:

$$\Delta I(x, y) = \frac{\partial^2 I(x, y)}{\partial x^2} + \frac{\partial^2 I(x, y)}{\partial y^2}$$

Image Laplacian

The kernel can be implemented as

0	0	0
1	-2	1
0	0	0

Second derivative kernel
(x direction)

0	1	0
0	-2	0
0	1	0

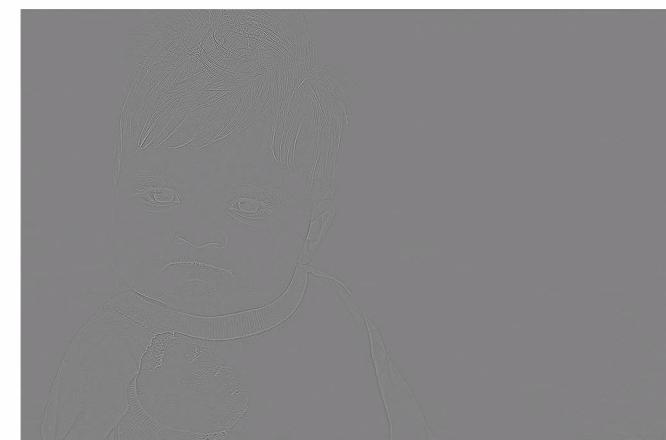
Second derivative kernel
(y direction)

0	1	0
1	-4	1
0	1	0

Laplacian kernel

```
K = fspecial('laplacian', 0);
J = imfilter(I, K);
imshow(J, []);
```

ΔI



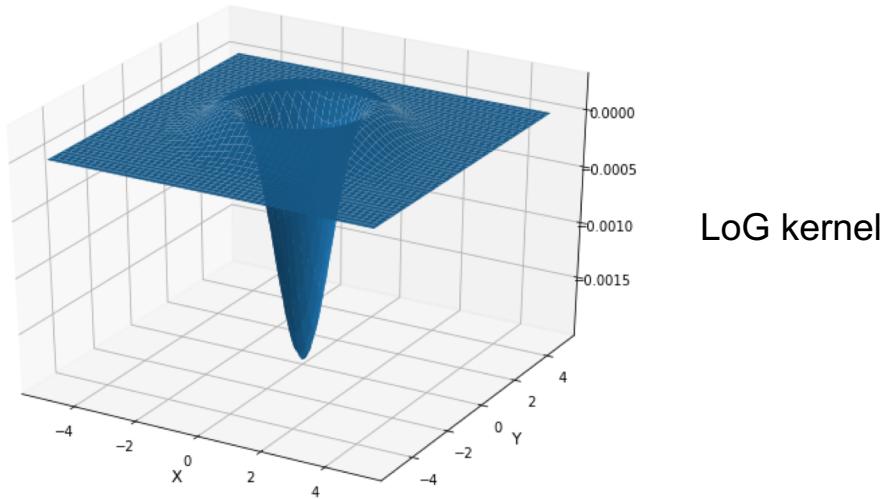
Laplacian of Gaussian

- The second derivative is more sensitive to noise than the first derivative
- We can also smooth the image using a Gaussian kernel, which results in a Laplacian of Gaussian filter:

$$K_1 = \frac{1}{2\pi\sigma^2} e^{\frac{-(x^2+y^2)}{2\sigma^2}}$$

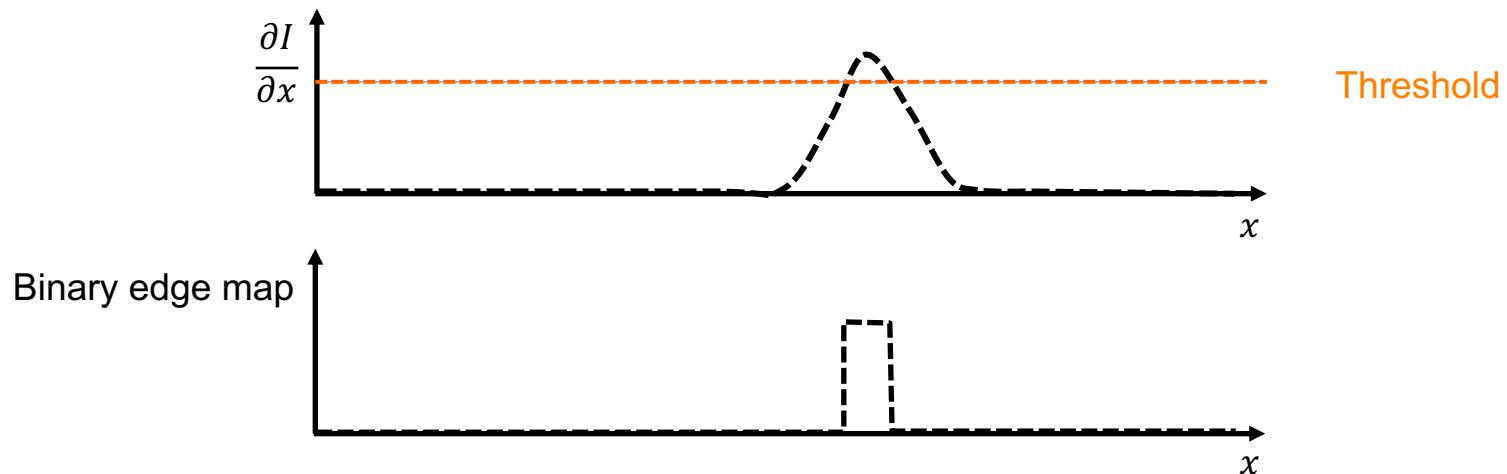
$$K_2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

- $K_3 = K_2 * K_1$ is the Laplacian of Gaussian (LoG) kernel
- Built-in Matlab function: `K = fspecial('log', hsize, sigma)`



Towards binary edge maps

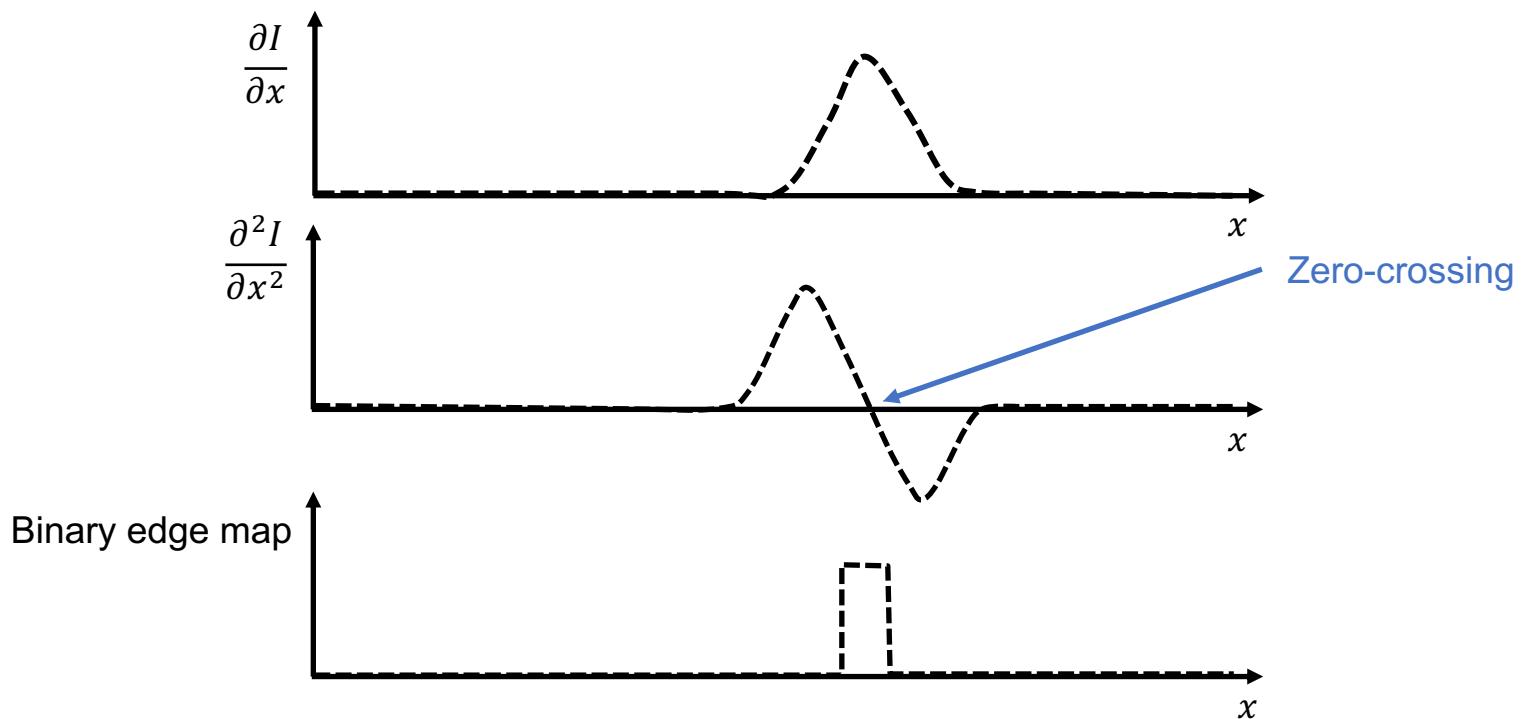
- In most applications, a binary edge map is required: 0s everywhere, and 1s in correspondence of an edge
- **Edge-based methods** look for the points of **maximum-response**: the higher the value of the gradient, the more likely it is for the pixel to belong to an edge
- To create a binary edge map, a **threshold** is usually used:



- This is not a straightforward process:
 - A low threshold will be susceptible to noise and detect irrelevant edges
 - A high threshold will miss subtle edges or result in fragmented ones

Towards binary edge maps

- Laplacian-based methods, on the other hand, look for **zero-crossings** to identify edges:



- Theoretically, there is no need to decide a threshold
- However, second derivatives are very sensitive to noise, so in the end some manual thresholds can still be needed

Canny edge detector

Canny [1] proposed the following criteria for an edge detector:

- **Good detection**

There should be a low probability of failing to mark real edge points, and low probability of falsely marking non-edge points

- **Good localisation**

The points marked as edge points by the operator should be as close as possible to the centre of the true edge

- **Single response**

Only one response to a single edge

Canny edge detector

Canny [1] proposed a gradient-based edge detection algorithm with the following main steps:

1. Perform Gaussian filtering to suppress noise
2. Calculate the gradient magnitude and direction
3. Apply non-maximum suppression (NMS) to get a single response for each edge
4. Perform hysteresis thresholding to find potential edges

Canny edge detector

1. Perform Gaussian filtering to suppress noise:

The choice for the parameter σ in the Gaussian kernel has a major impact on the result of the detector:

- High $\sigma \rightarrow$ detect large scale structures
- Low $\sigma \rightarrow$ detect fine details

2. Calculate the gradient magnitude and direction

These two steps can be done in two ways:

- In two separate steps: Gaussian filter + Prewitt or Sobel filter
- In one step using the derivative of Gaussian filters

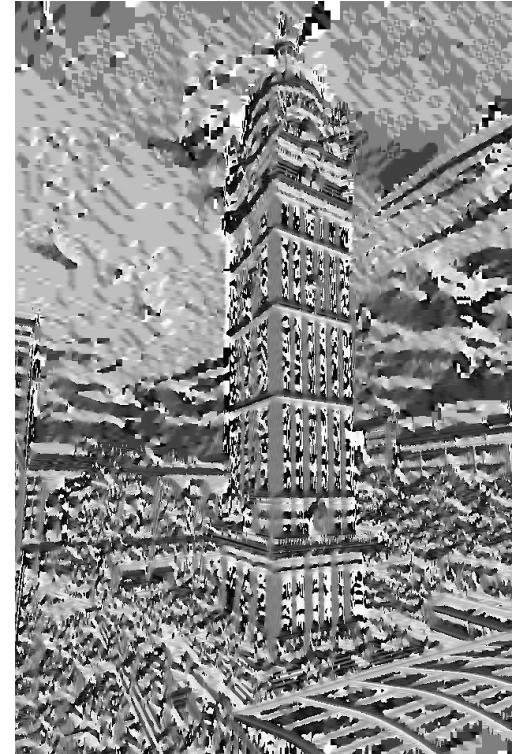
Canny edge detector



Input image



Gradient magnitude



Gradient direction

Canny edge detector

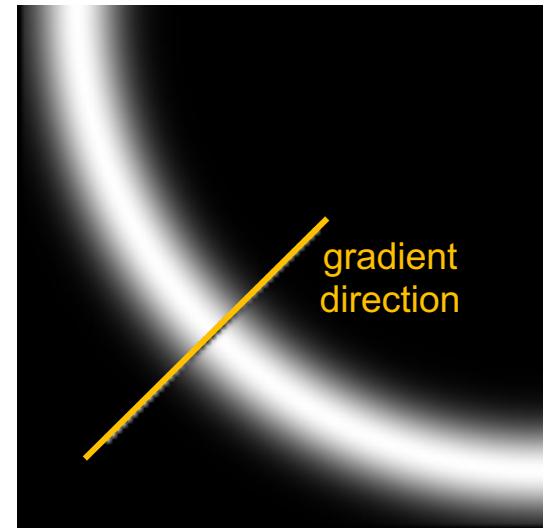
3. Apply non-maximum suppression (NMS) to get a single response for each edge:

NMS is an edge thinning technique: suppress all the gradient values (by setting them to 0) except the local maxima, which indicate locations with the sharpest change of intensity value

Steps for NMS:

- For each pixel q , check whether the gradient magnitude is a local maximum along gradient direction
- If it is, leave magnitude untouched; if it is not, set magnitude to 0

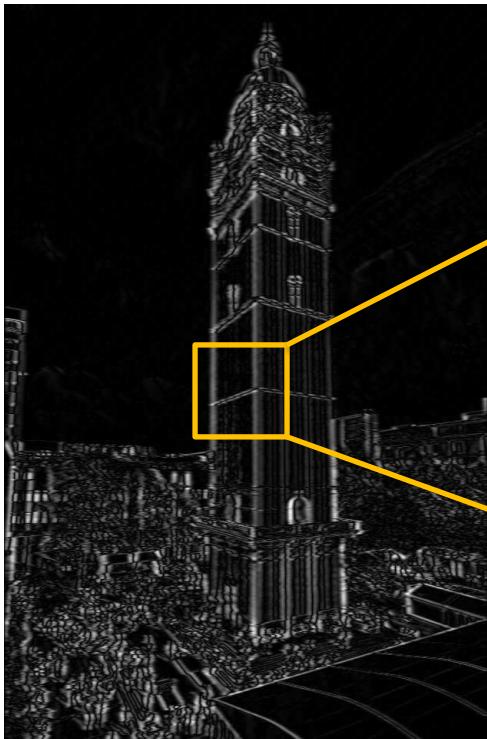
Usually, the gradient directions are quantized into a set of discrete ones:
horizontal, vertical, diagonal and anti-diagonal



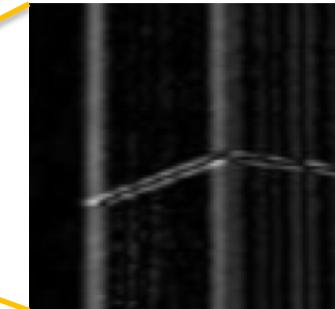
Gradient magnitude

Canny edge detector

3. Apply non-maximum suppression (NMS) to get a single response for each edge:



Before NMS



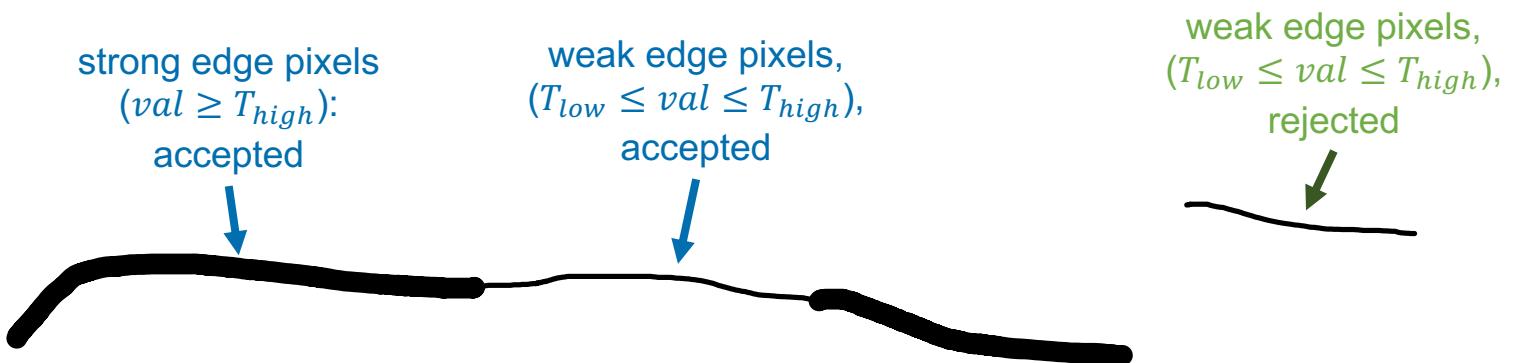
After NMS



Canny edge detector

4. Perform hysteresis thresholding to find potential edges

- Hysteresis is the lagging of an effect, an inertia
- Hysteresis thresholding requires the definition of two thresholds: T_{high} , T_{low}
- In edge detection, hysteresis thresholding means that
 - If a pixel gradient magnitude is larger than T_{high} , it is accepted as an edge
 - If a pixel gradient magnitude is lower than T_{low} , it is rejected
 - If in between, it will be accepted if it is connected to an edge



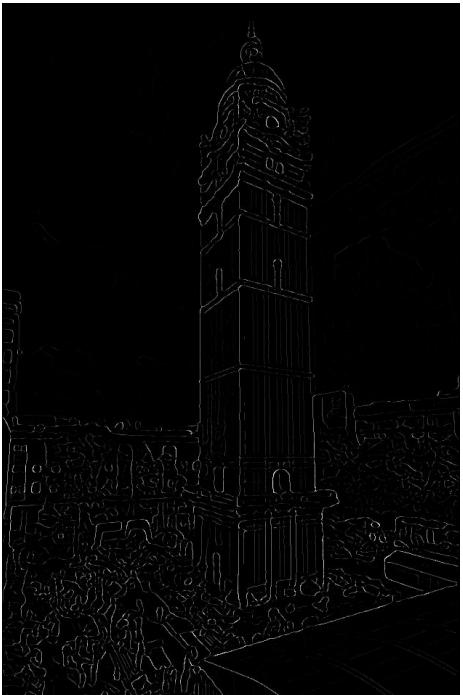
Canny edge detector



Input image



Gradient magnitude



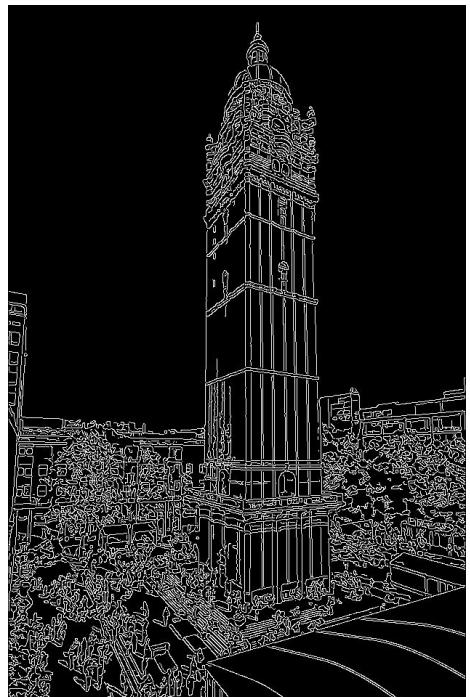
NMS



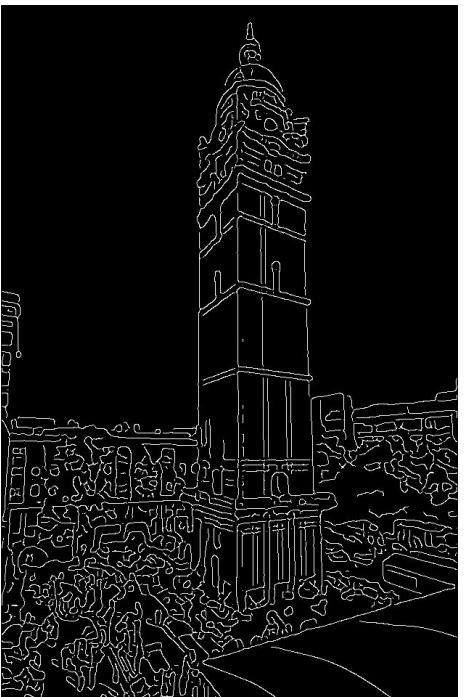
Hysteresis thresholding

Canny edge detector

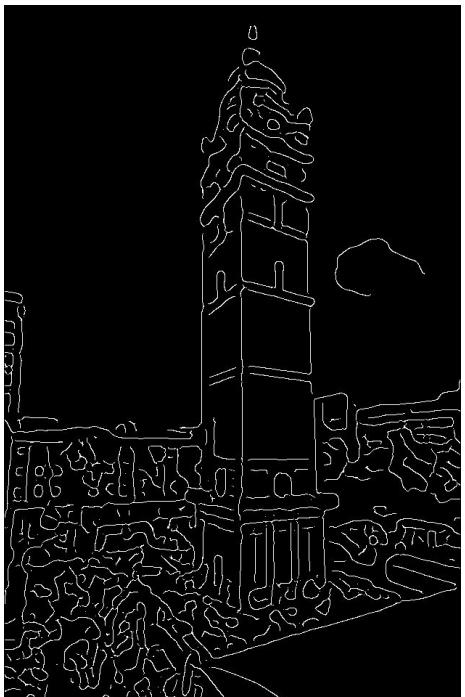
Effect of the Gaussian kernel parameter:



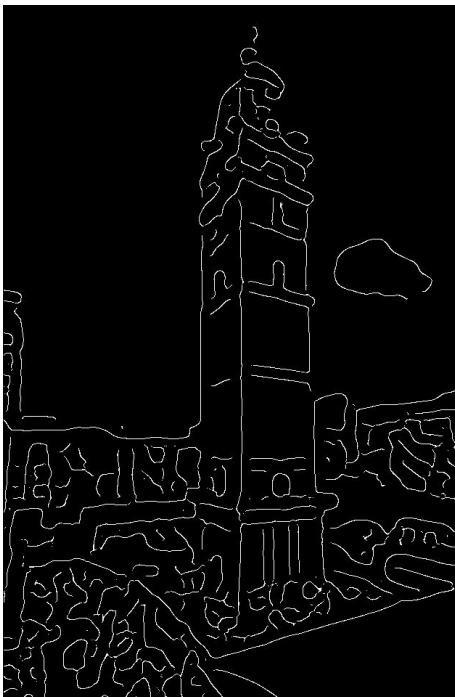
$\sigma = 1$



$\sigma = 3$



$\sigma = 5$



$\sigma = 7$

- High $\sigma \rightarrow$ detect large scale structures
- Low $\sigma \rightarrow$ detect fine details

Additional resources

All of the presented edge detection strategies are implemented in a single function in Matlab: `edge`

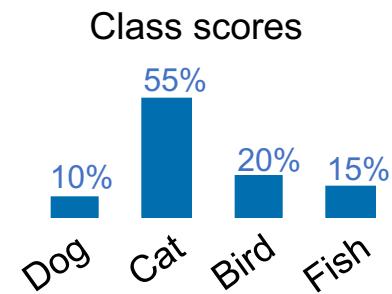
Common CV tasks

Image classification

- Identify the class of the object shown in the image
- Often (but not necessarily) based on class scores estimation



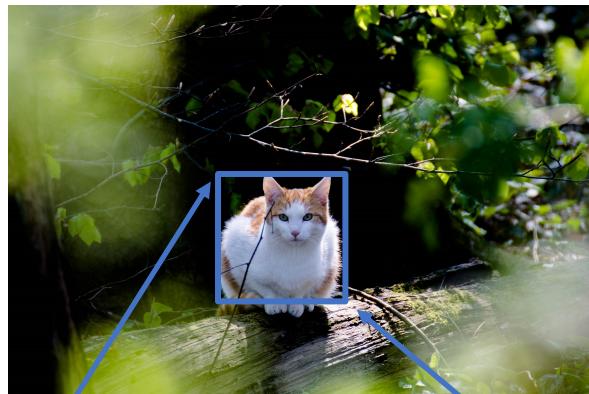
Class label: "Cat"



Common CV tasks

Image classification & localization

- Identify the class and the location of the object (regression of bounding box coordinates)
- Assume only one object per image



- Class label: “Cat”
- Bounding box coordinates: (x, y, w, h)

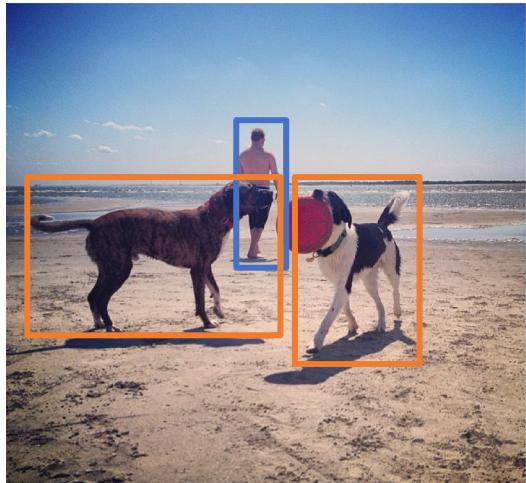
(x, y)

$(x + w, y + h)$

Common CV tasks

Object detection

- Identify the class and the location of the objects in the image
- Assume **more than one object** per image



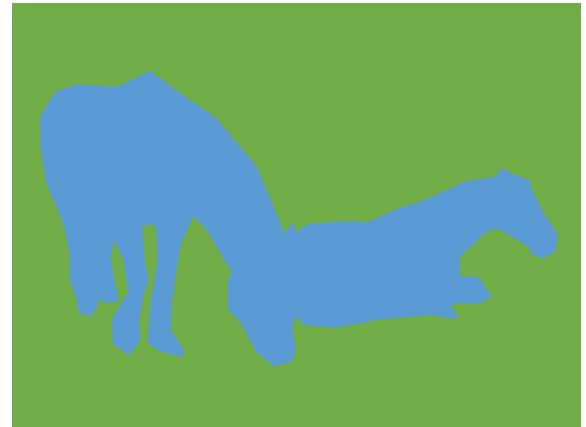
For each object in the image:

- Class label: “Dog”
- Bounding box coordinates: (x, y, w, h)

Common CV tasks

Semantic segmentation

- Assign a label to each pixel



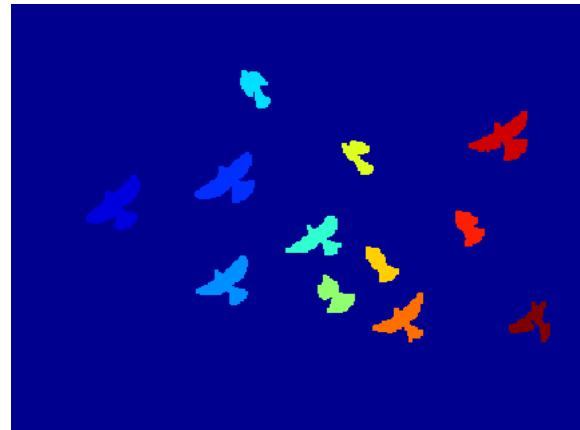
Common CV tasks

Instance segmentation

- Assign a label to each pixel
- **Distinguish** between different **instances**
- Combination of object detection and semantic segmentation



Instance segm.



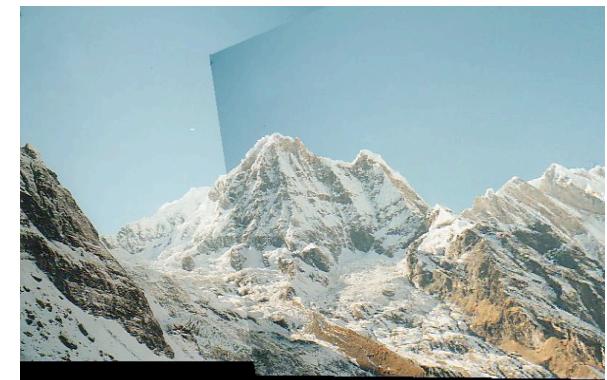
Common CV tasks

Image matching/alignment/registration

- Identify the (rigid? affine? non-rigid?) transformation which aligns two or more images
- Used for multi-modal comparison, image stitching, etc.



Matching



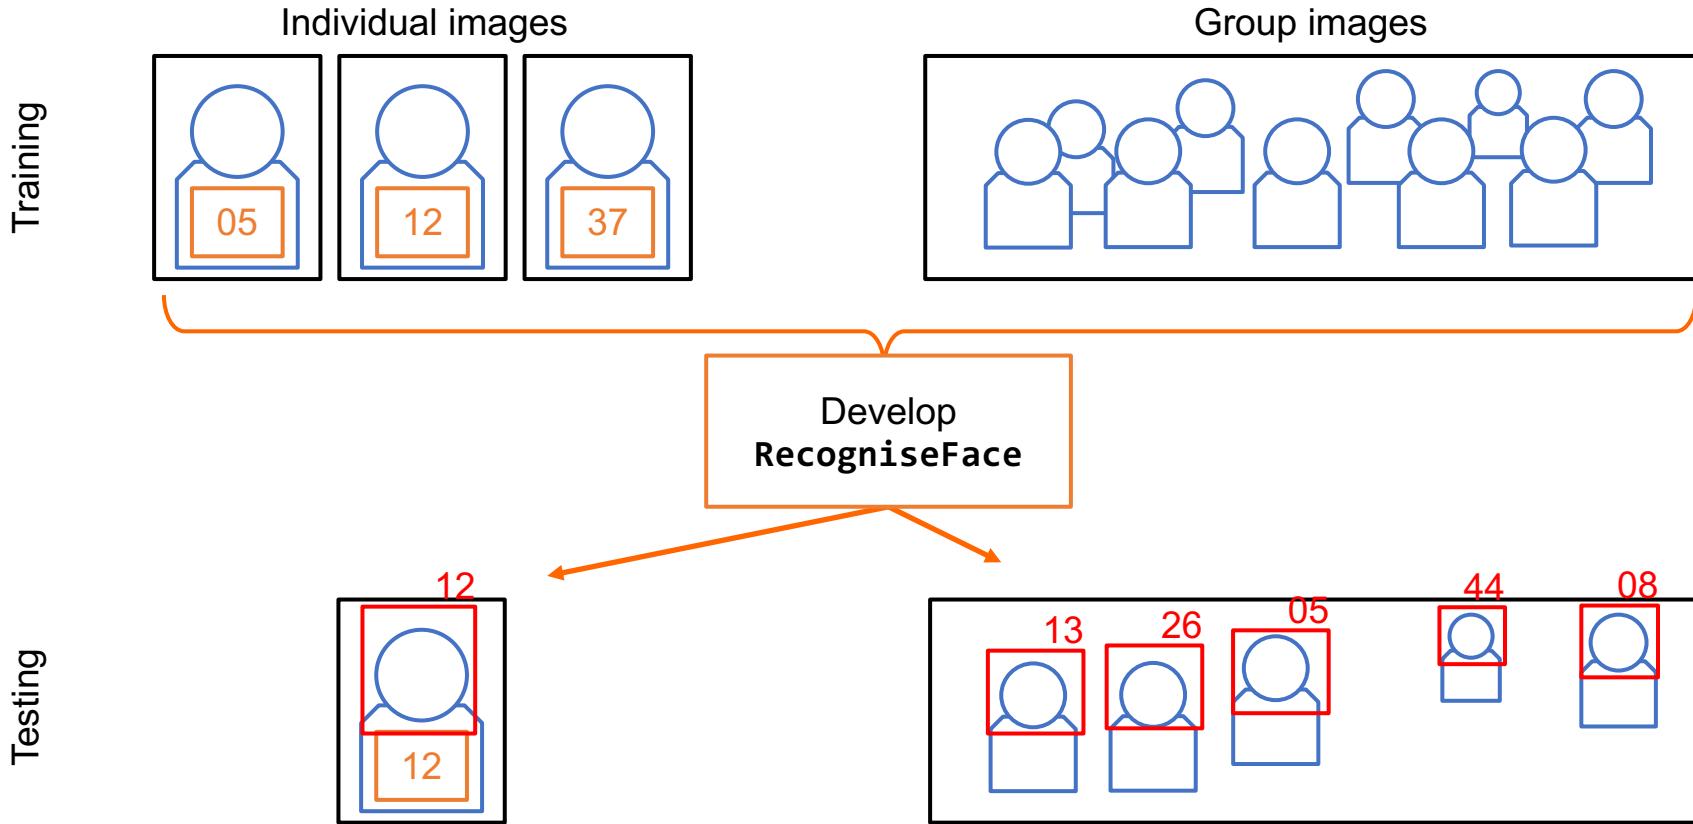
Overview of next week's lecture

Image segmentation:

- Intensity-based methods
 - Thresholding
 - K-means clustering
- Watershed
- Active contours
 - Snakes
 - Level-sets
- Graph cuts
- Superpixels
- Quantifying results: how to tell if we achieved a good segmentation

Coursework: face recognition

- Deadline (UGs & PGs): **Thursday 30 April 2020, 5pm**
- Coding: **Matlab** (strongly recommended). Python solutions will also be accepted (but support from academic staff limited to Matlab)
- Develop a function **RecogniseFace** that identifies faces and assigns IDs



Coursework: face recognition

- **RecogniseFace** must include different solutions, i.e. combinations features-classifiers
- The training and evaluation procedures will be carried out by you, starting from the data split
- Short videos (i.e. “Live Photos”) have been added too: extract frames from them if you believe you need additional training/evaluation data
- At least for the best implemented combination of features-classifier, provide performance metrics (e.g. % of accurately classified faces in a portion of the evaluation set)
- Your **RecogniseFace** will be tested on unseen individual/group images (not available on the folders on Moodle)
- **Additional task for PGs:** if prompted, allow for **RecogniseFace** to modify the detected faces (e.g. cartoonify, add content e.g. beard, etc.)

Coursework: face recognition

- You will need to submit:
 - **RecogniseFace** and other developed code (commented and referenced!)
 - Readme.txt to run your code
 - Trained models
 - Report describing your approach and showing results
 - Short (**max 3 mins**) screen capture video showing your function running on images from your evaluation set
- You **must not** submit the training/evaluation dataset
- Package all the listed items in a zip and submit to Moodle. Given the 200MB limit, the **models alone** can be uploaded on personal cloud space (e.g. Dropbox) and **linked in the readme.txt file**
- **Your models will not be retrained.** Your **RecogniseFace** function must run as is, so package it properly and test it before submission
- Don't freak out if you don't get very high accuracy in some cases!